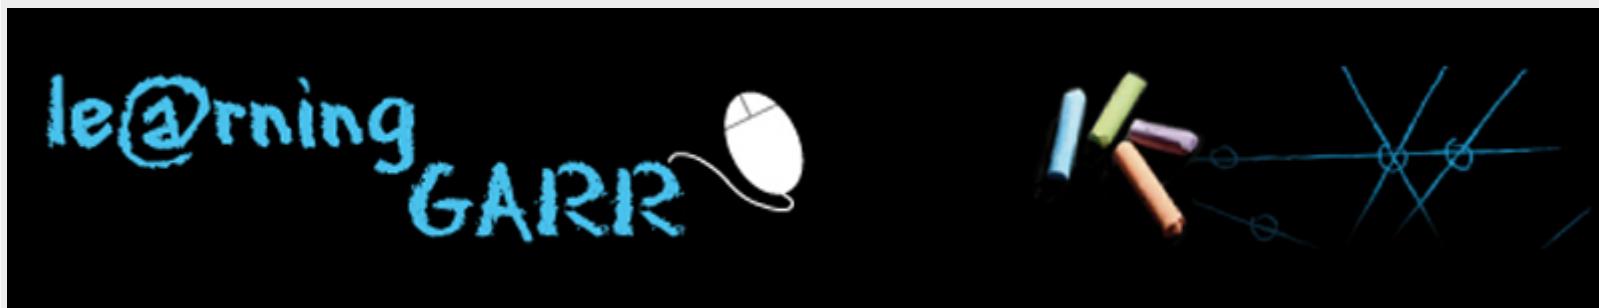


DOCKER VS VIRTUALIZZAZIONI

28 MAGGIO 2018

ROMA

RIGRAZIAMENTI



ABOUT ME

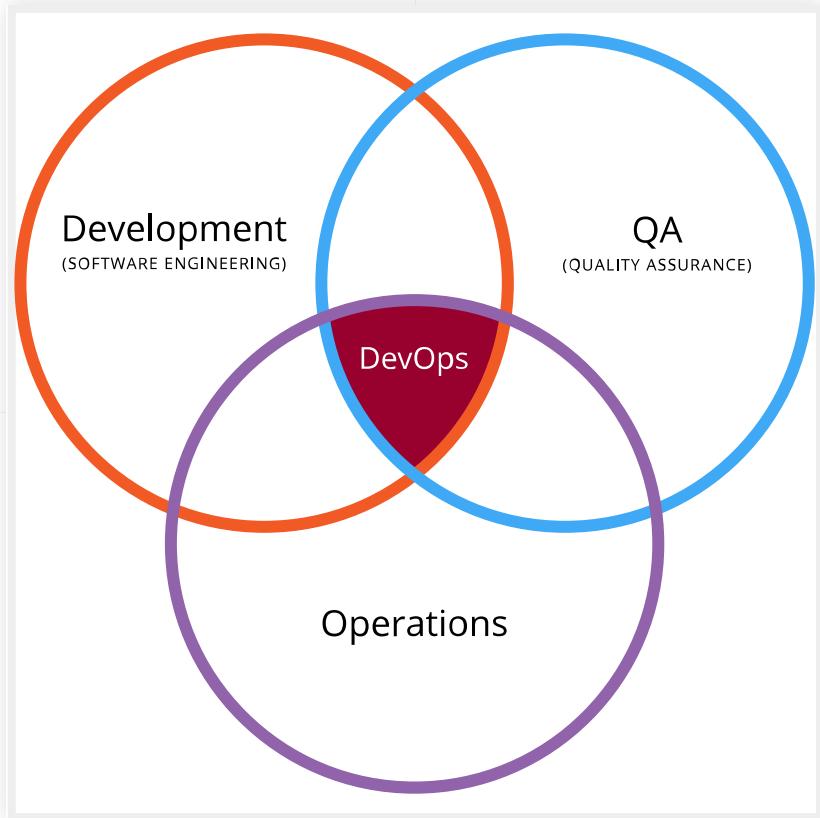
- ITPro > 25 year
- Senior Software Developer at INNOTEAM s.r.l.
- Linux Foundation Member
- CTS Senior Member at Federation IDEM AAI
- ICT Administrator at Università Politecnica delle Marche
- Istruttore Informatico at Comune di Chiaravalle (AN)
- Member of WindowServer.it community
- Member of DevMarche community

I.I.I.D.

**INFRASTRUTTURA, ISOLAMENTO,
IDEMPOTENZA E DEVOPS**

**I VANTAGGI E I VINCOLI DELLA DEFINIZIONE FORMALE DI
UN'INFRASTRUTTURA**

DEVOPS



Una completa automazione assicura la facile ripetibilità dei rilasci e riduce gli errori nell'operazione.

INFRASTRUTTURA

Insieme di materiale ed elementi singoli che collaborano a sostenere, produrre ed erogare beni e servizi. Nel mondo ICT sono *hardware* e/o *artefatti software* che collaborano per risolvere problemi complessi.

ISOLAMENTO

Il concetto di *Isolamento* dipende dalla definizione di due domini che chiameremo per convenzione *Interno* ed *Esterno* e dalla definizione di una superficie di separazione tra Interno ed Esterno chiamata per convenzione *Confine*. L'*Isolamento* si applica sul *Confine* e condiziona lo scambio tra *Interno* ed *Esterno*.

IDEMPOTENZA

È la proprietà di un'operazione che applicata una o più volte non cambia il risultato prodotto.

L'applicazione multipla ed asincrona della funzione con questa proprietà non introduce collisioni o comportamenti non deterministici.

STATO

Contenuto informativo temporaneo o persistente nell'artefatto che annulli la proprietà di *idempotenza*. Per ripristinare la proprietà di *idempotenza*, l'artefatto delega ad elementi esterni la gestione dei propri stati.

DEFINIZIONE FORMARE D'INFRASTRUTTURA

Descrizione oggettiva dell'*infrastruttura* la cui implementazione viene delegata ad una procedura automatica.

ESEMPIO DI DEFINIZIONE FORMALE

```
version: '2'

services:
  web:
    build: ./php/
    volumes:
      - /var/www:/var/www/html:rw
    ports:
      - "80:80"
      - "443:443"
    links:
      - mysql
  environment:
    XDEBUG_CONFIG: remote_host=10.211.55.2
```

VANTAGGI

- Resilienza
- Scalabilità Dinamica
- Documentazione
- Replicabilità
- Diaccoppiamento dall'Hardware
- Test d'Infrastruttura automatizzabili
- Condivisione

VINCOLI

- Esatta conoscenza dei domini che concorrono nell'Infrastruttura
- Oneroso per aggiustamenti di lieve entità
- Rigore nel Workflow di lavoro
- Studio, Analisi, Test, Tempo

ask who? discover
where? how? why? challenge who?

QUESTIONS

ask who? discover
when? knowing clues
investigation ? why? ask
what known investigation

DOCKER

**UN FACILITATORE PER UTILIZZARE I
CONTAINER E LE TECNOLOGIE AD ESSI
COLLEGATE**

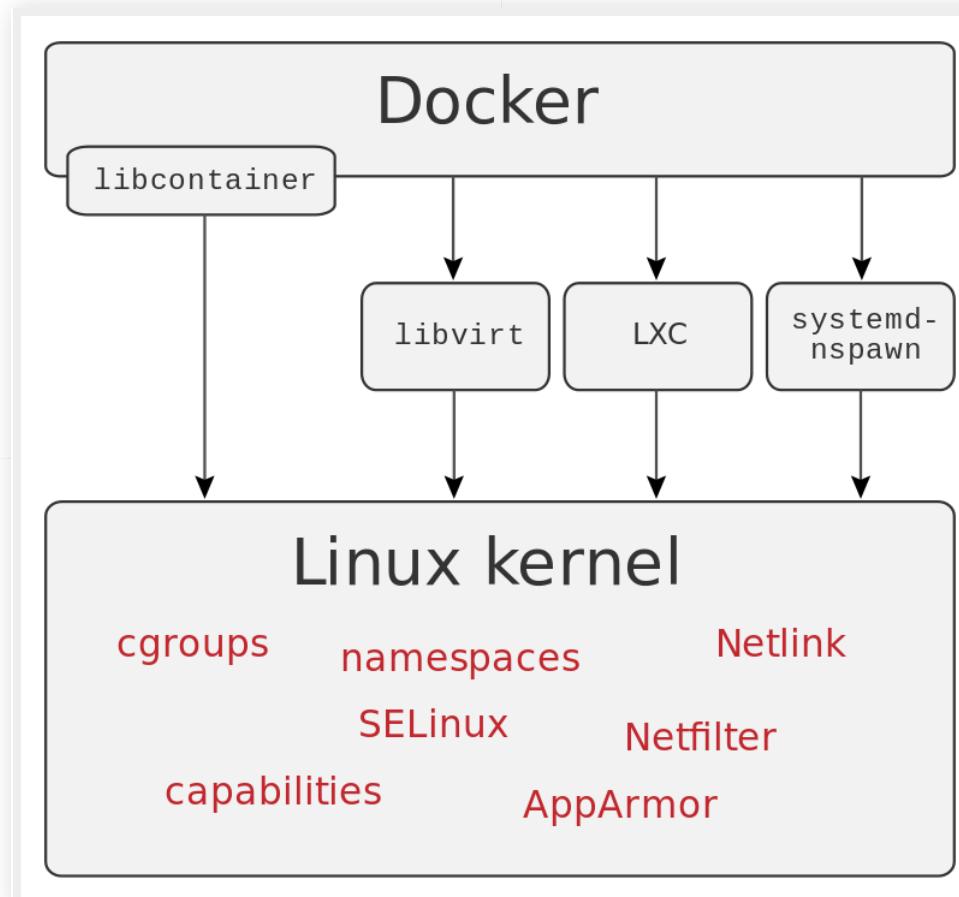
COSA È DOCKER

UN FACILITATORE?

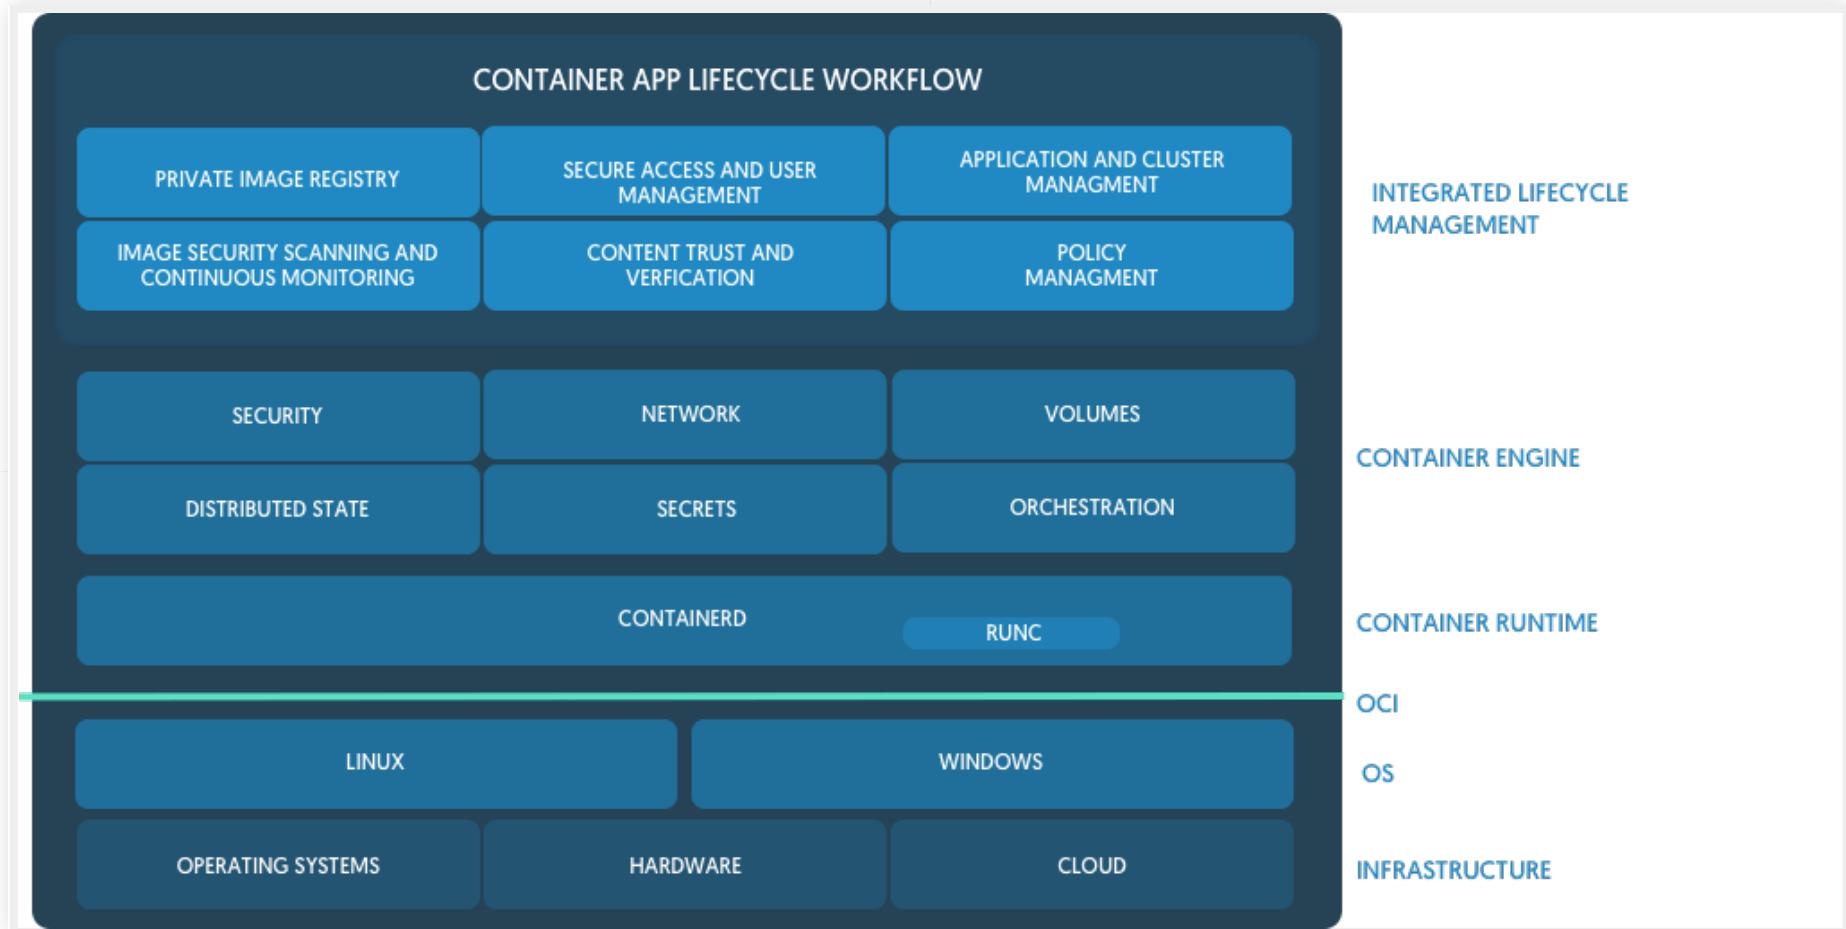
UN HUB DI TECNOLOGIE?

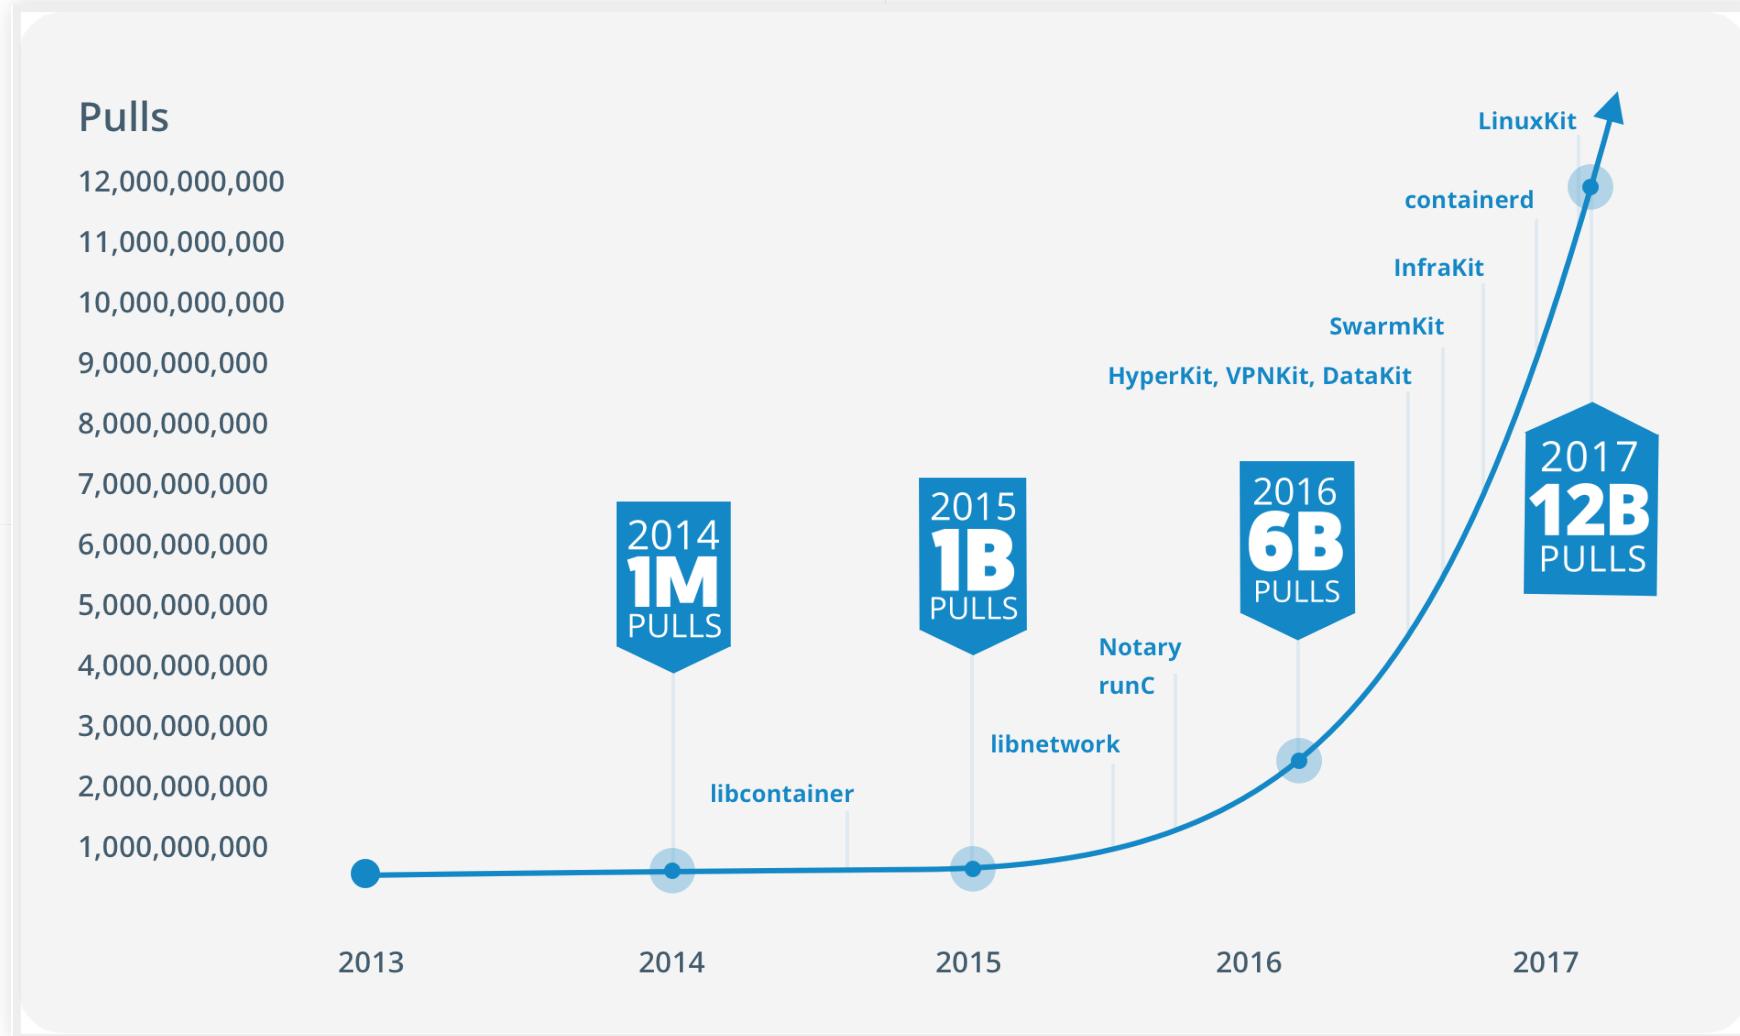
**UN SIGNIFICATO NUOVO ALLA PAROLA *INFRASTRUTTURA*
E UN MODO NUOVO PER COSTRUIRLA.**

DOCKER, GLI INIZI

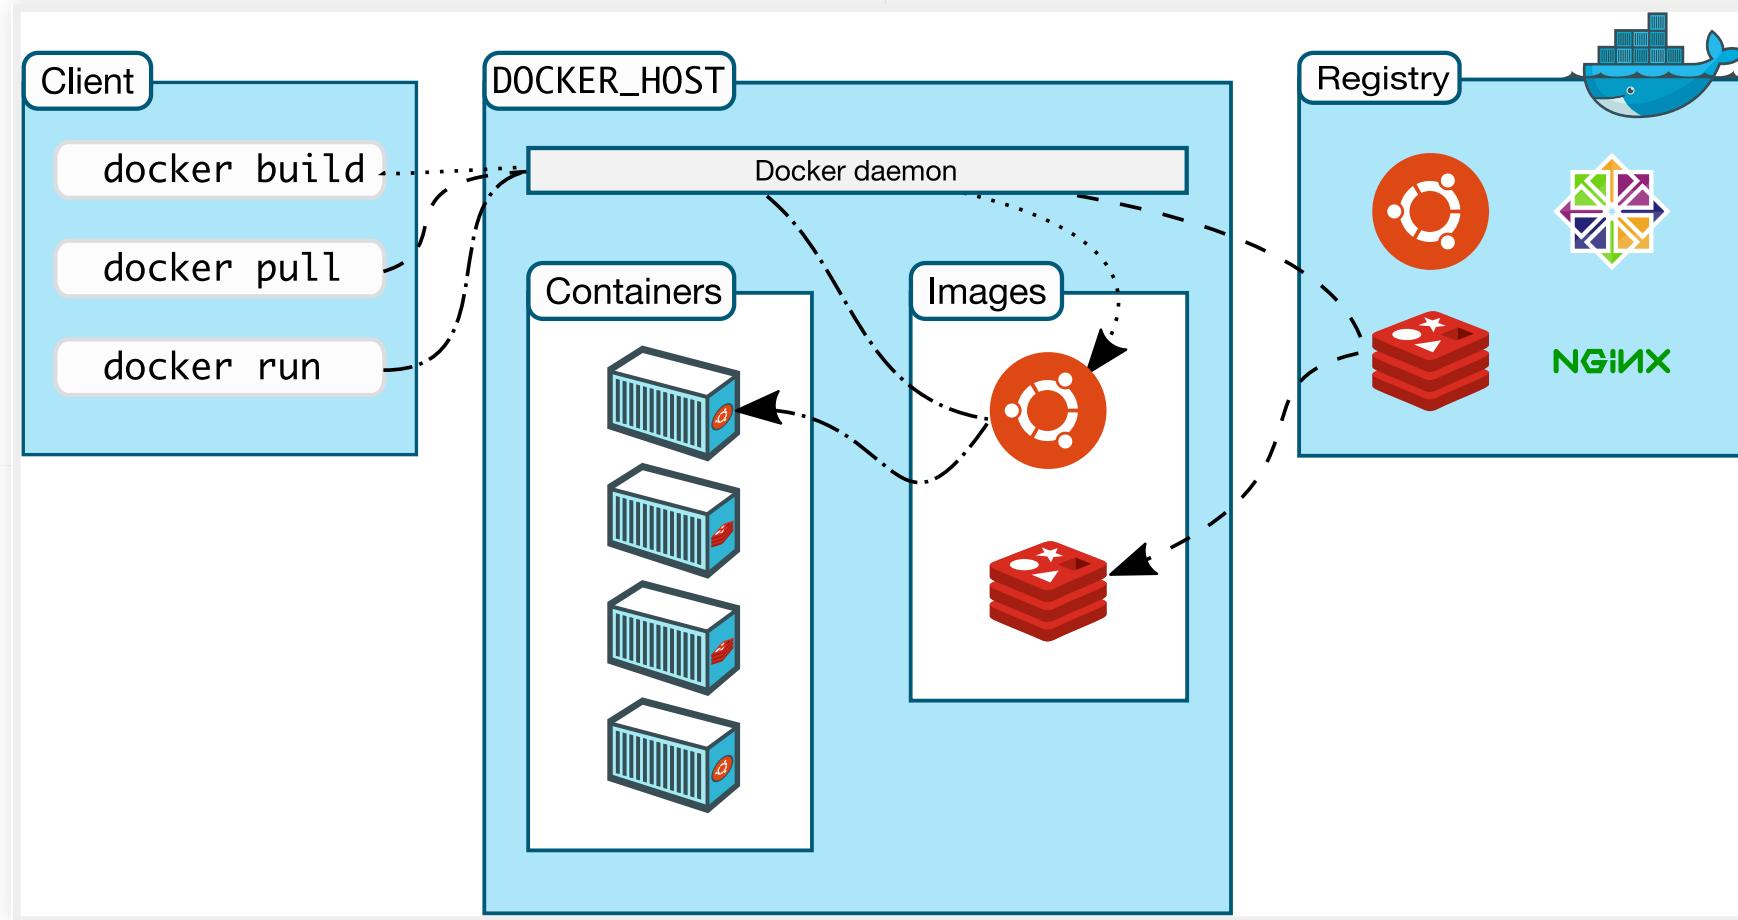


DOCKER OGGI





CONTAINERS WORKFLOW



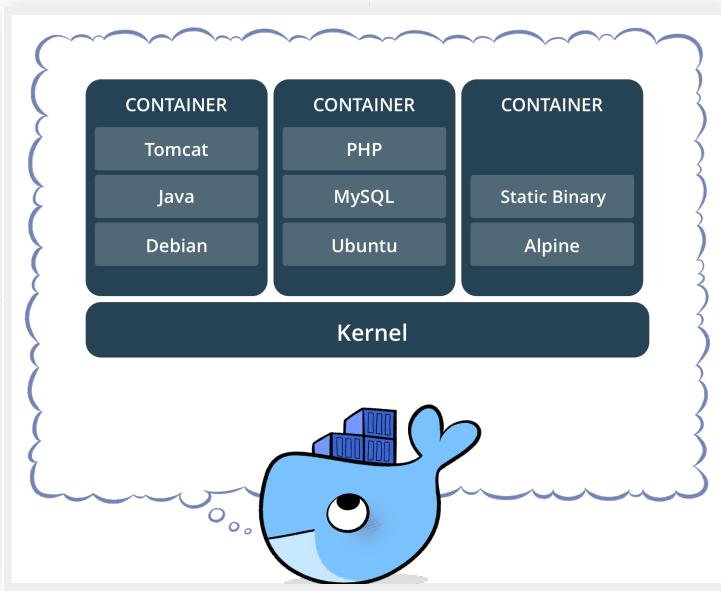
L'IMMAGINE SECONDO DOCKER.INC

A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings.

– [Official Docker website](#)

È un **artefatto**. Il **Container** è la sua rappresentazione operativa come il **processo** lo è del **codice**.

IL CONTAINER



Il *Container* è un elemento isolato e autosufficiente, accessibile attraverso la rete. Esso astrae a livello applicativo le dipendenze necessarie al codice.

DOCKER 101: INTRODUCTION TO DOCKER

Docker 101: Introduction to Docker and Containers (copy)

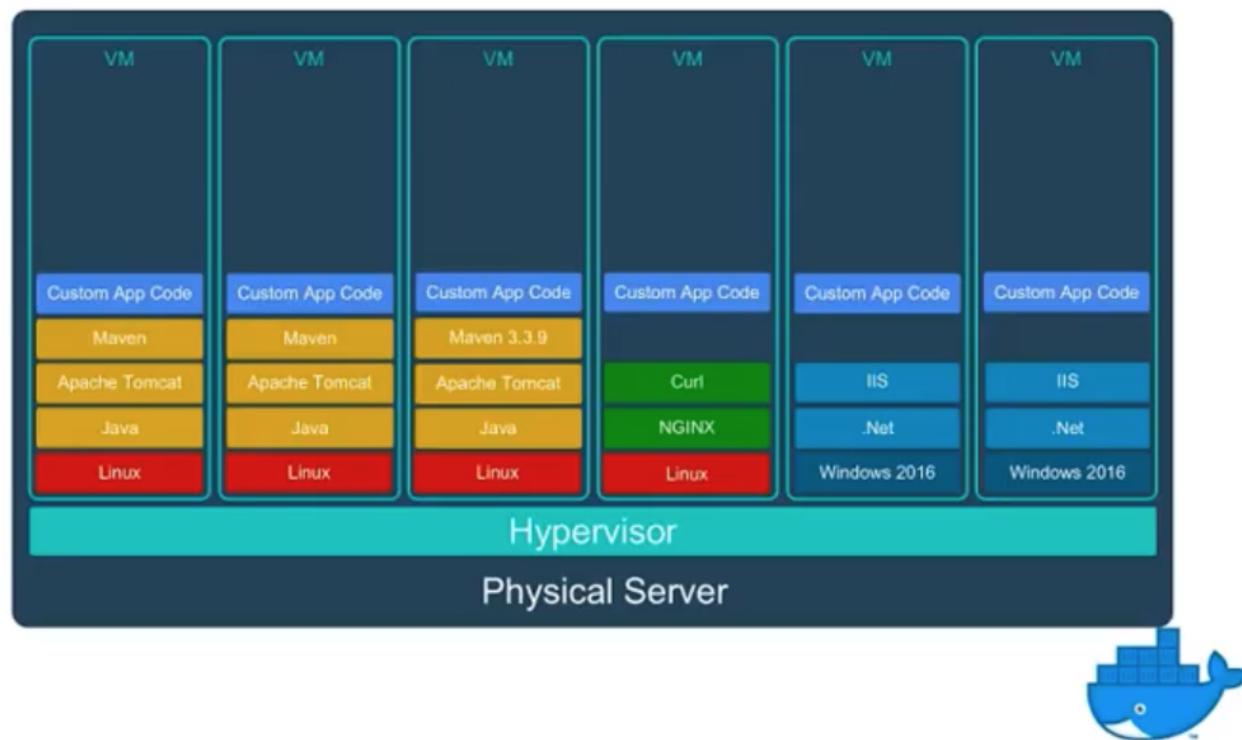
2.008 visualizzazioni



VM CONSOLIDATION

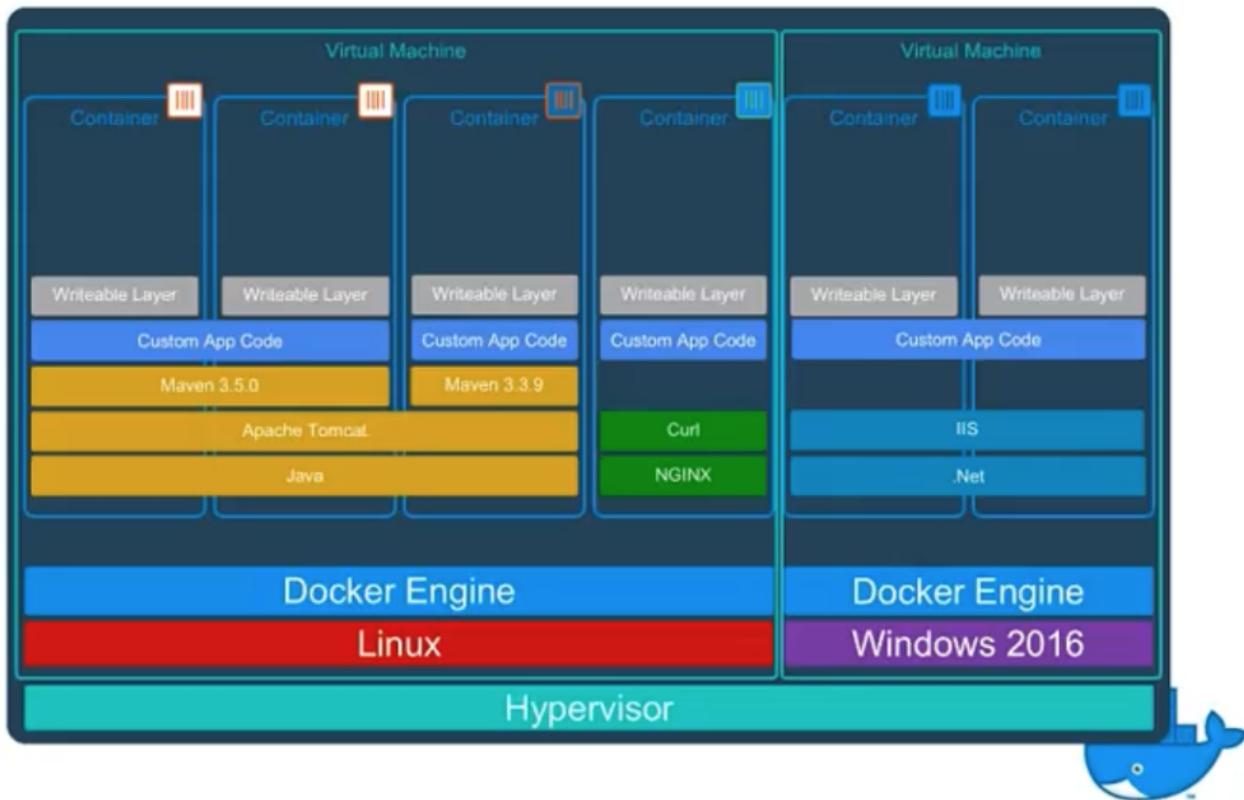
Our Current State: Virtual Machines

- One app per VM
- Separate OS copy
- Separate app framework
- "Black box" to hypervisor
- App is tracked elsewhere



CONTAINER CONSOLIDATION

Docker Delivers Savings Through Consolidation



CREAZIONE DI UN CONTAINER

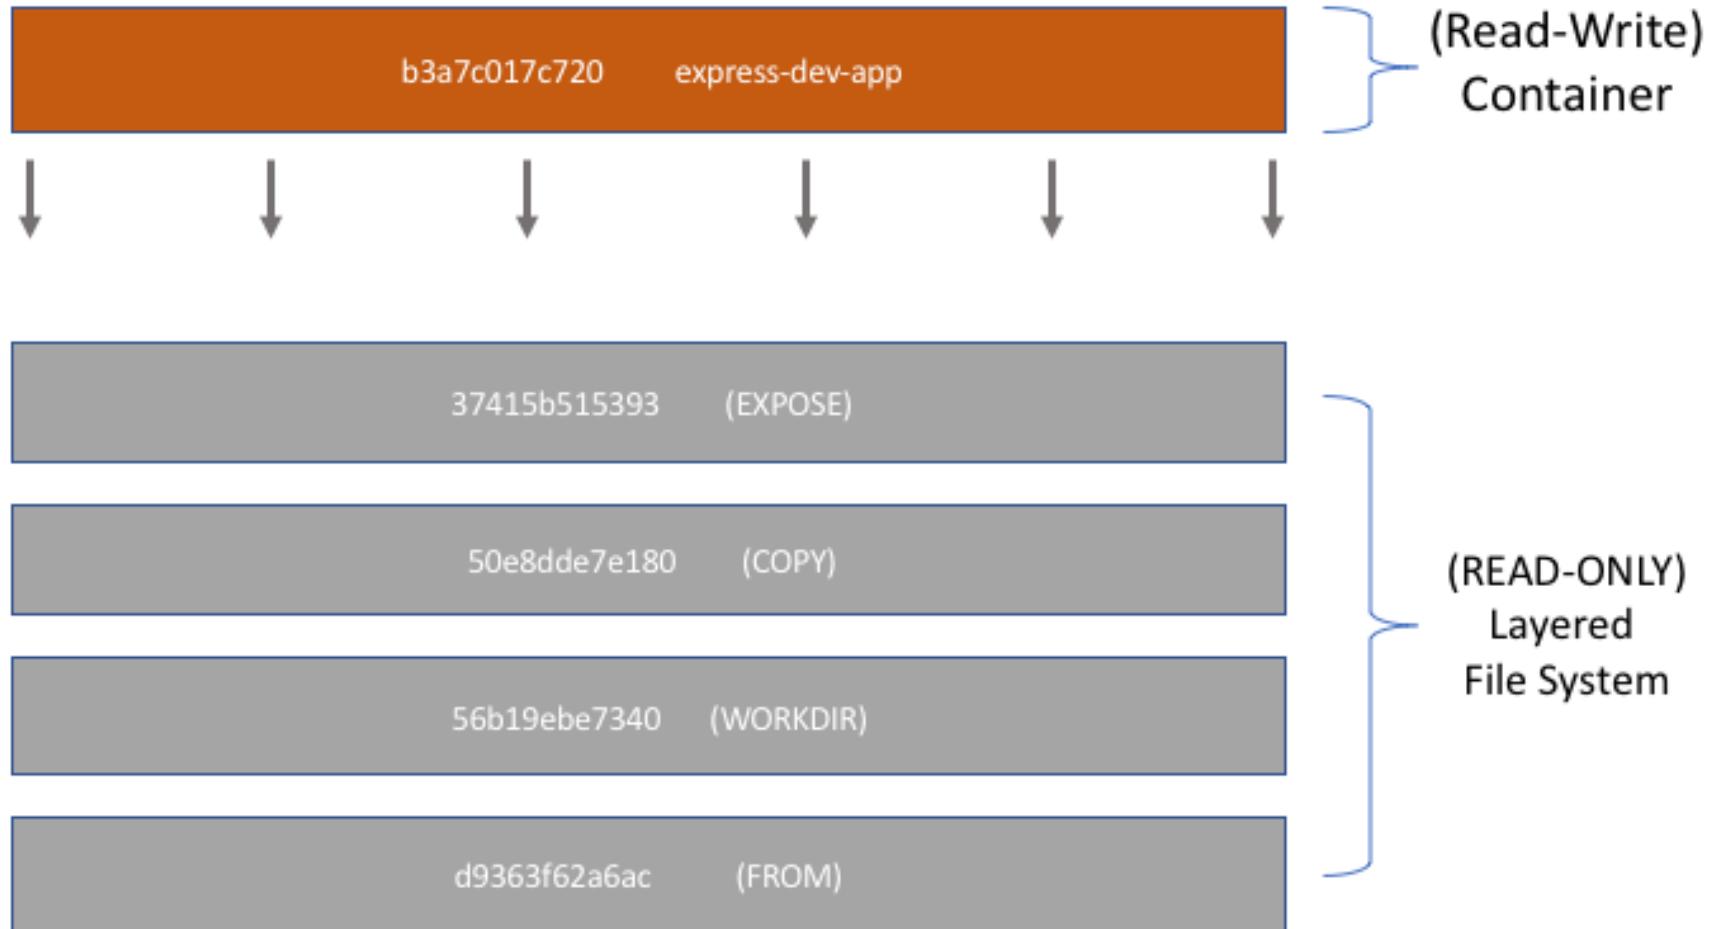
```
bash-4.4$ docker run -it ubuntu bash
Unable to find image 'ubuntu:latest' locally
latest: Pulling from library/ubuntu
1be7f2b886e8: Pull complete
6fbc4a21b806: Pull complete
c71a6f8e1378: Pull complete
4be3072e5a37: Pull complete
06c6d2f59700: Pull complete
Digest: sha256:e27e9d7f7f28d67aa9e2d7540bdc2b33254b452ee8e60f3888
Status: Downloaded newer image for ubuntu:latest
root@66170252cf62:/#uname -a
Linux 66170252cf62 4.9.75-linuxkit-aufs #1 SMP Tue Jan 9 10:58:17
```

CREAZIONE DI UN'IMMAGINE

```
FROM php:5.6-apache
# Install XDebug
RUN yes | pecl install xdebug \
    && echo "zend_extension=$(find /usr/local/lib/php/extensions/ \
    && echo "xdebug.remote_enable=on" >> /usr/local/etc/php/conf.d \
    && echo "xdebug.remote_autostart=off" >> /usr/local/etc/php/c

# Install PHP exrensions.
RUN apt-get update \
    && apt-get install -y \
        imagemagick \
        graphicsmagick \
        zip \
        unzip \
        wget \
        ....\
```

From: ./02_Docker/code/01/dockerfile



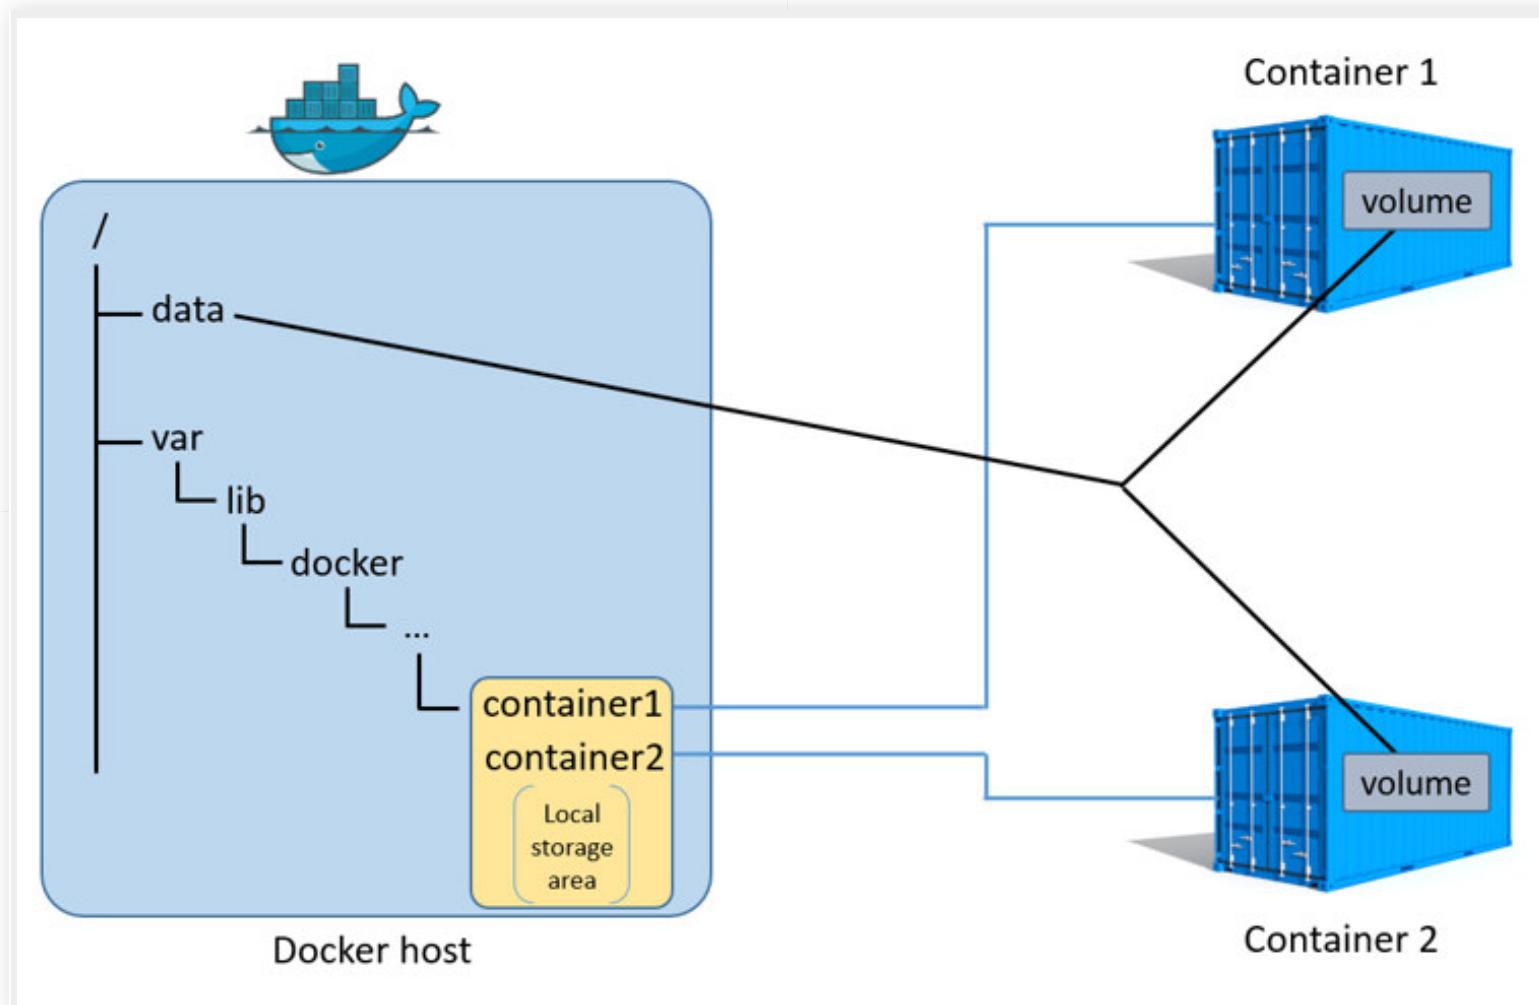
COSA SUCCIDE DOPO UN *DOCKER RUN*



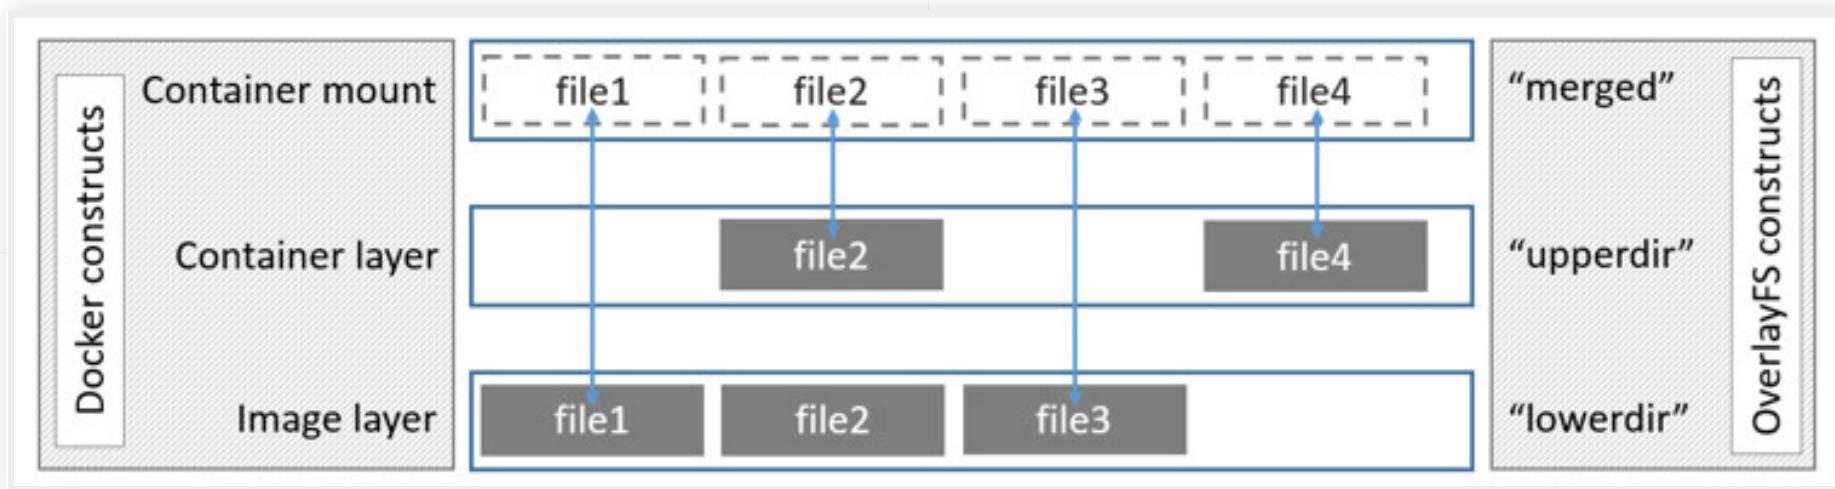
LAYERS CONDIVISI



HOST-FILESYSTEM CONDIVISI



FILES VISION IN OVERLAYFS



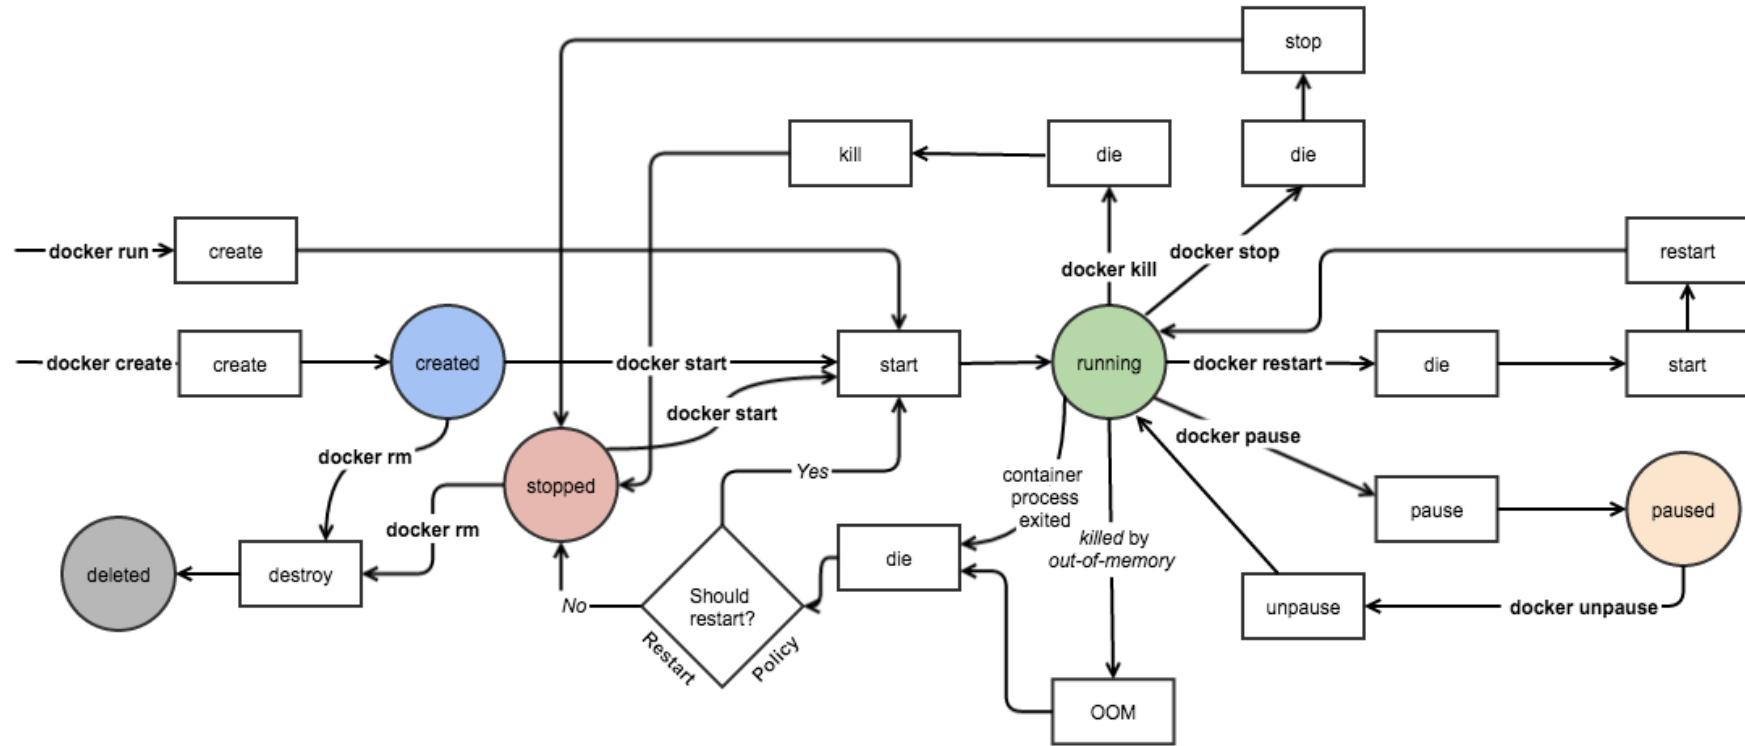
GLOSSARIO DOCKER BASE

- *Docker Engine*: deamon gestore dei elementi di un *Docker Host*
- *Docker Client*: CommandLine Interface per interagire con la *Docker Engine* locale o remota
- *Docker Image*: readonly template generato da `docker build` da cui avviare uno o più Container
- *Docker Container*: artefatto attivo prodotto da `docker run` a partire da un'immagine

GLOSSARIO DOCKER AVANZATO

- *Docker Host*: Sistema fisico o virtuale in cui viene eseguita la *Docker Engine*
- *Docker Registry*: Repository di immagini Pubblico hub.docker.com o Privato
- *Docker Compose*: Orchestrator di infrastrutture multicontainer organizzati in servizi
- *Docker Swarm*: Cluster multihost ad orchestrazione automatica

WORKFLOW COMANDI DOCKER



- Build Image (Interpretazione del Dockerfile)

```
docker build dockerID/imagename:tag .
```

- Run Container (creazione del Container da un'immagine)

```
docker run -d dockerID/imagename:tag
```

- Pull Image (scarico di un'immagine)

```
docker pull dockerID/imagename:tag
```

- Push Image (invio di un immagine al repository)

```
docker push dockerID/imagename:tag
```


ANALISI DI UN DOCKERFILE

```
# vim:set ft=dockerfile:
FROM debian:jessie

# add our user and group first to make sure their IDs get assigned
RUN groupadd -r mysql && useradd -r -g mysql mysql

# add gosu for easy step-down from root
ENV GOSU_VERSION 1.10
RUN set -ex; \
    \
    fetchDeps=' \
        ca-certificates \
        wget \
    '; \
    apt-get update; \
    \
    curl -L https://github.com/tianon/gosu/releases/download/v1.10/gosu-debian-1.10 \
        | sed -e 's@http://apt-get:https://apt-get@gosu@g' \
        | tee /usr/local/bin/gosu \
        | chmod +x /usr/local/bin/gosu
```

From: ./02_Docker/code/02/dockerfile

ask who? discover
where? how? why? challenge who?

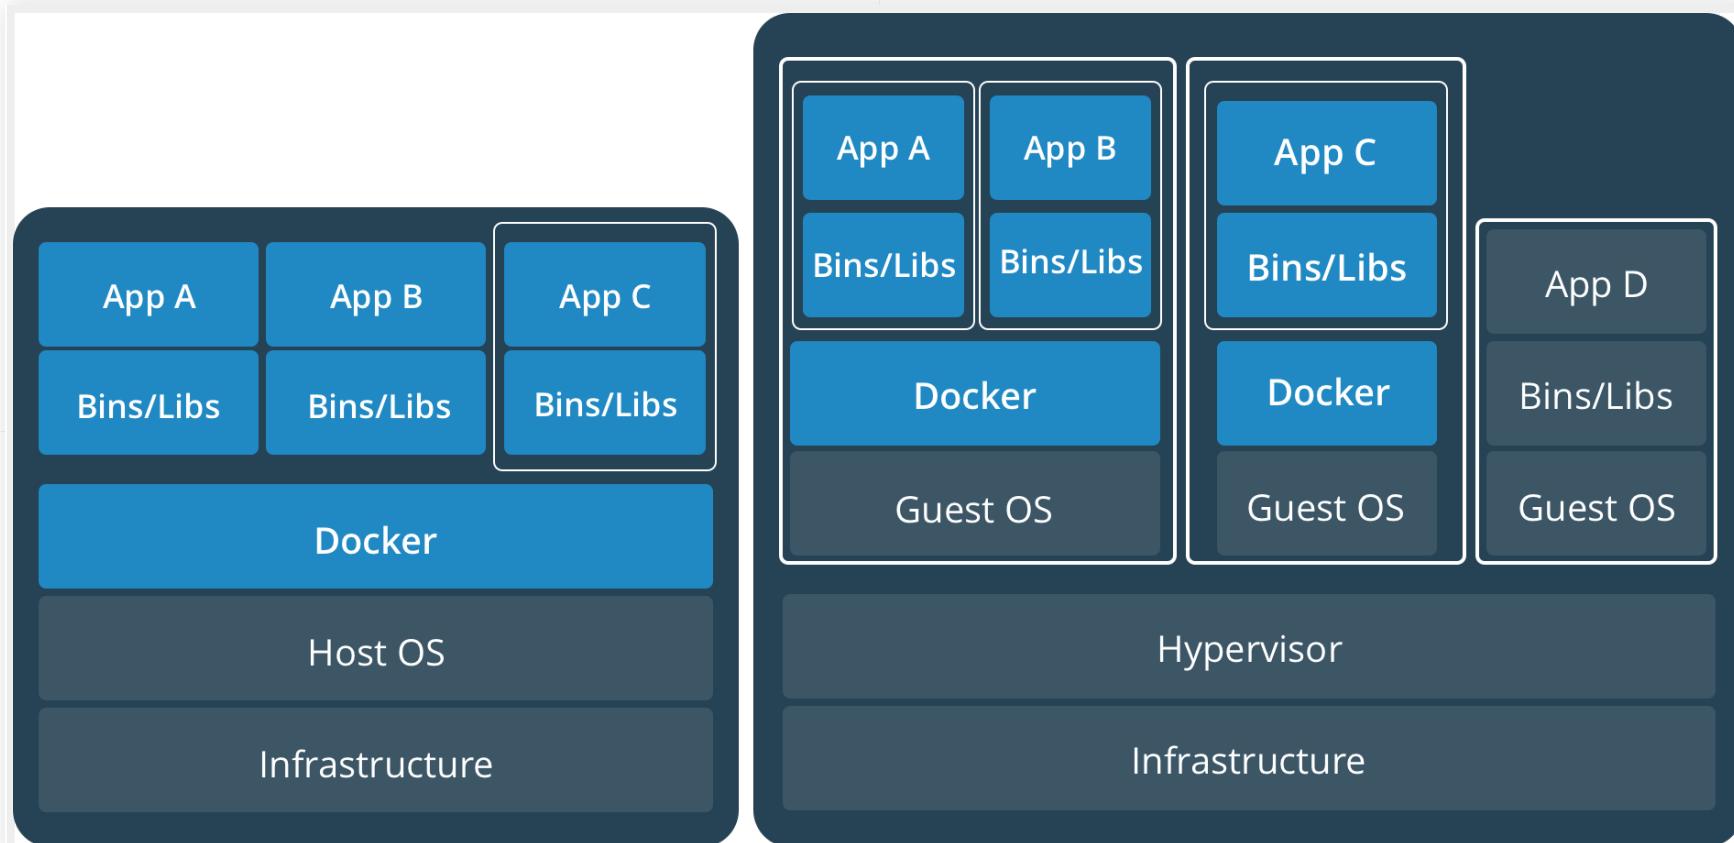
QUESTIONS

ask who? discover
when? knowing clues why? known
investigation ? ask
what?

CONTAINER VS HYPERVISOR

**LE DUE FACCE DELLA STESSA MEDAGLIA:
PRO, CONTRO ED ESEMPI D'USO**

VM E CONTAINER A CONFRONTO



QUANDO USO LE VM

- Hardware da simulare
 - Es. *MAME is a multi-purpose emulation framework.*
- Software monolitico e fortemente interconnesso
 - Es. *GUI unita alla Business logic*
- Software GUI e Console esclusi quelli Web-based
- Software non scalabile
- Quando lo sforzo di Automazione delle Ops produce un ROI negativo

QUANDO USO I CONTAINER

- Boundaries dei domini ben definiti
- Software cli-based su OS Linux o Windows
- Software su architettura Client/Server connessi attraverso Network TCP/IP
- Alta scalabilità e automazione dei processi Op
- Assenza di stato nei Container Applicativi, gli stati sono persistiti da altri elementi architetturali

DEPLOY AUTOMATIZZATO DI INFRASTRUTTURE A VM

- Definizione di una *GoldenImage* o *Template* come iso base d'installazione
- Deploy automatico e personalizzazioni
 - scripting • Puppet
 - juju • Chef
 - Vagrant • Mesosphere DC/OS
 - Ansible • Kubernetes

DEPLOY AUTOMATIZZATO DI INFRASTRUTTURE A CONTAINER

SCRIVO UN *DOCKER-COMPOSE.YML*

Docker, *out-of-the-box* si occupa di implementarlo.

DA VMA CONTAINER

**RICHIEDE UN PO' DI ESERCIZIO ED
UN DIVERSO APPROCCIO MENTALE.**

- No Gestione Manuale!!! -> Lavoro che andrà perso!
- No Gestione Network!!! -> L'infrastruttura è autonoma in questo!
- No Configurazioni Arlecchino!!! -> Creano *Single-Point-Of-Failure!*
- No Know-how Silos!!! -> Quando serve il *Master-of-Know-how* non si trova!

ask who? discover
where? how? why? challenge who?

QUESTIONS

ask who? discover
when? knowing clues
investigation ? why? ask
what known investigation

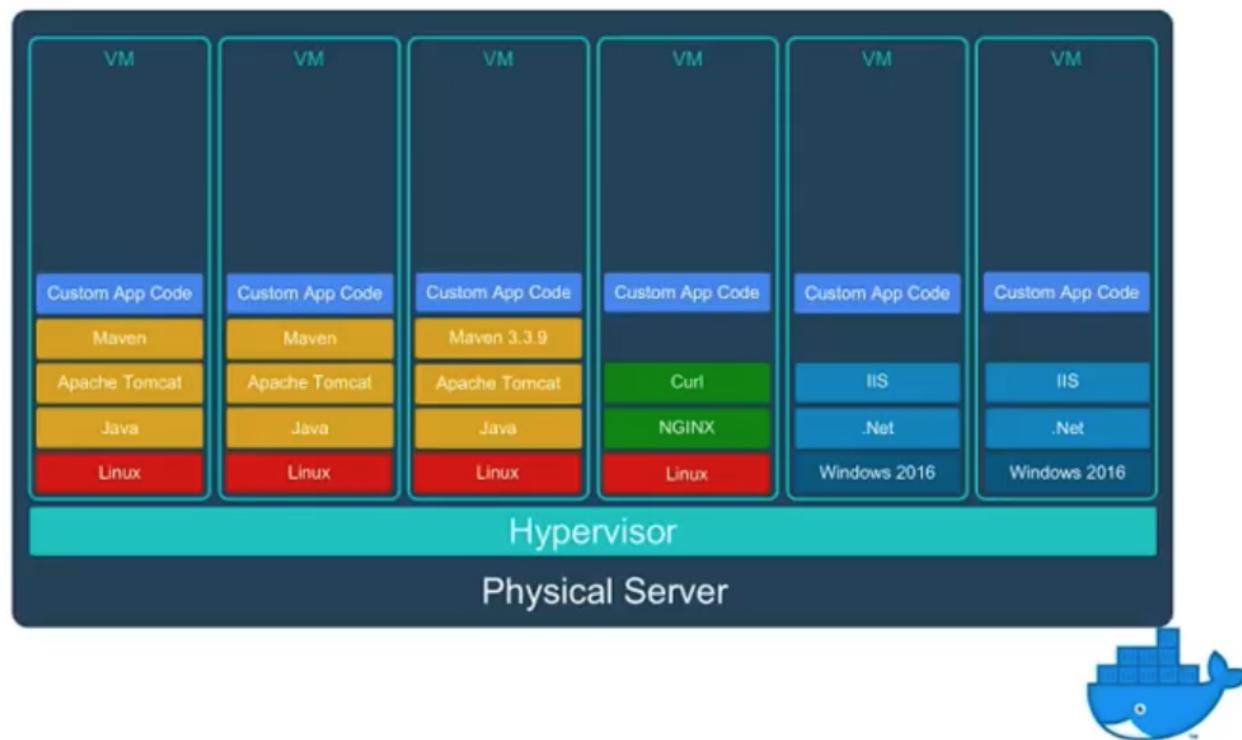
DALLA VM AI SERVIZI

**UN MODO PER RIMESCOLARE LE CARTE A
NOSTRO VANTAGGIO**

ARCHITETTURA A VM

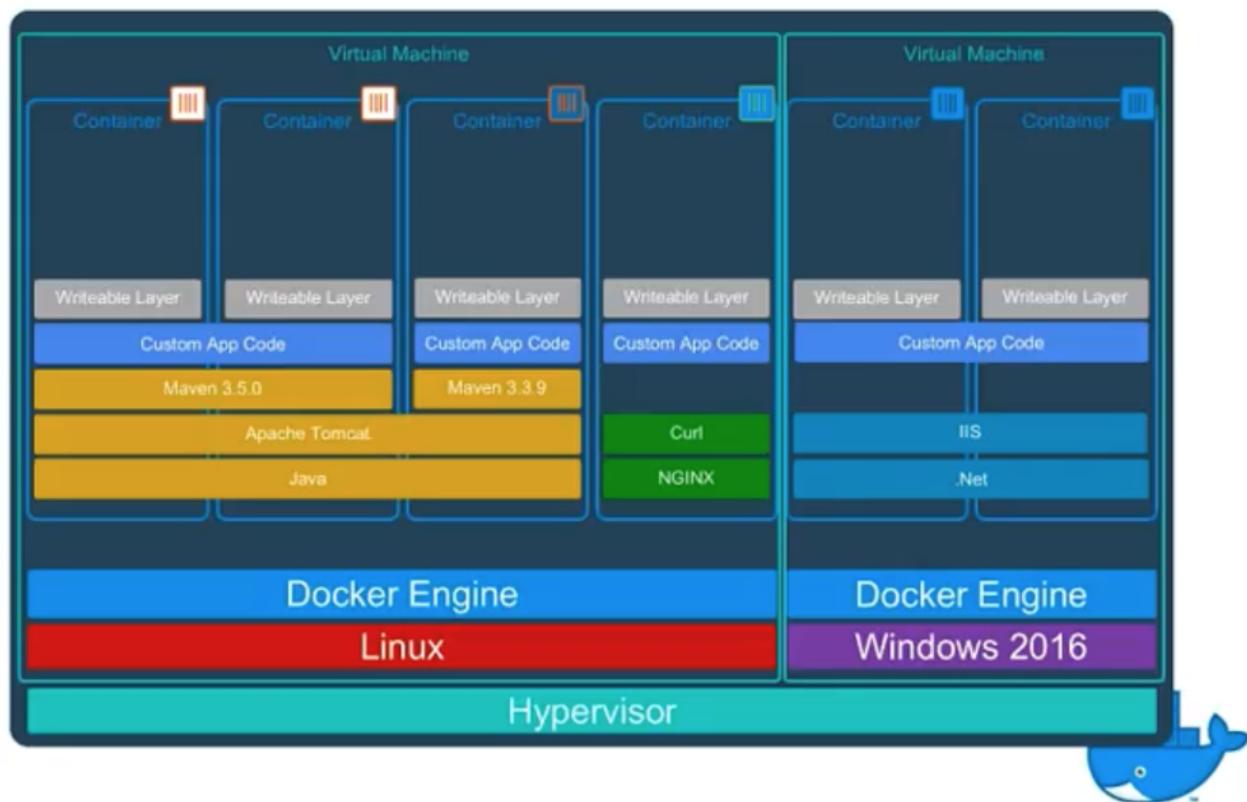
Our Current State: Virtual Machines

- One app per VM
- Separate OS copy
- Separate app framework
- "Black box" to hypervisor
- App is tracked elsewhere



ARCHITETTURA A CONTAINER

Docker Delivers Savings Through Consolidation



ANALISI DI UN *DOCKER-COMPOSE.YML*

```
version: '3'
services:
  db:
    image: mysql
    restart: always
    hostname: db
    domainname: develop.local
    container_name: db
    environment:
      MYSQL_ROOT_PASSWORD: secretadmin
    ports:
      - "3316:3306"
    volumes:
      - dbvolume01:/var/lib/mysql
    mail:
```

From: ./04_VM2Servizi/code/01/dockerfile

ask who? discover
where? how? why? challenge who?

QUESTIONS

ask who? discover
when? knowing clues
investigation ? why? ask
what known investigation

SVILUPPO SOFTWARE

ESEMPIO D'USO DI DOCKER NELLO SVILUPPO DI UNA PROCEDURA PHP

OBIETTIVO: CONTAINER COME VIRTUALENV PER PHP

Non installare la versione dell'interprete php
necessario allo sviluppo, ma usare l'immagine PHP
ufficiale per avere un Container interattivo nel proprio
sistema.

PASSI DA SEGUIRE

- Analizzare la configurabilità dell'immagine ufficiale di *PHP* su *DockerHub*
- Analizzare come la propria IDE s'interfaccia all'interprete *PHP* e a *Docker*
- Con una serie di test e prove arrivare alla propria *Comfort-Zone*

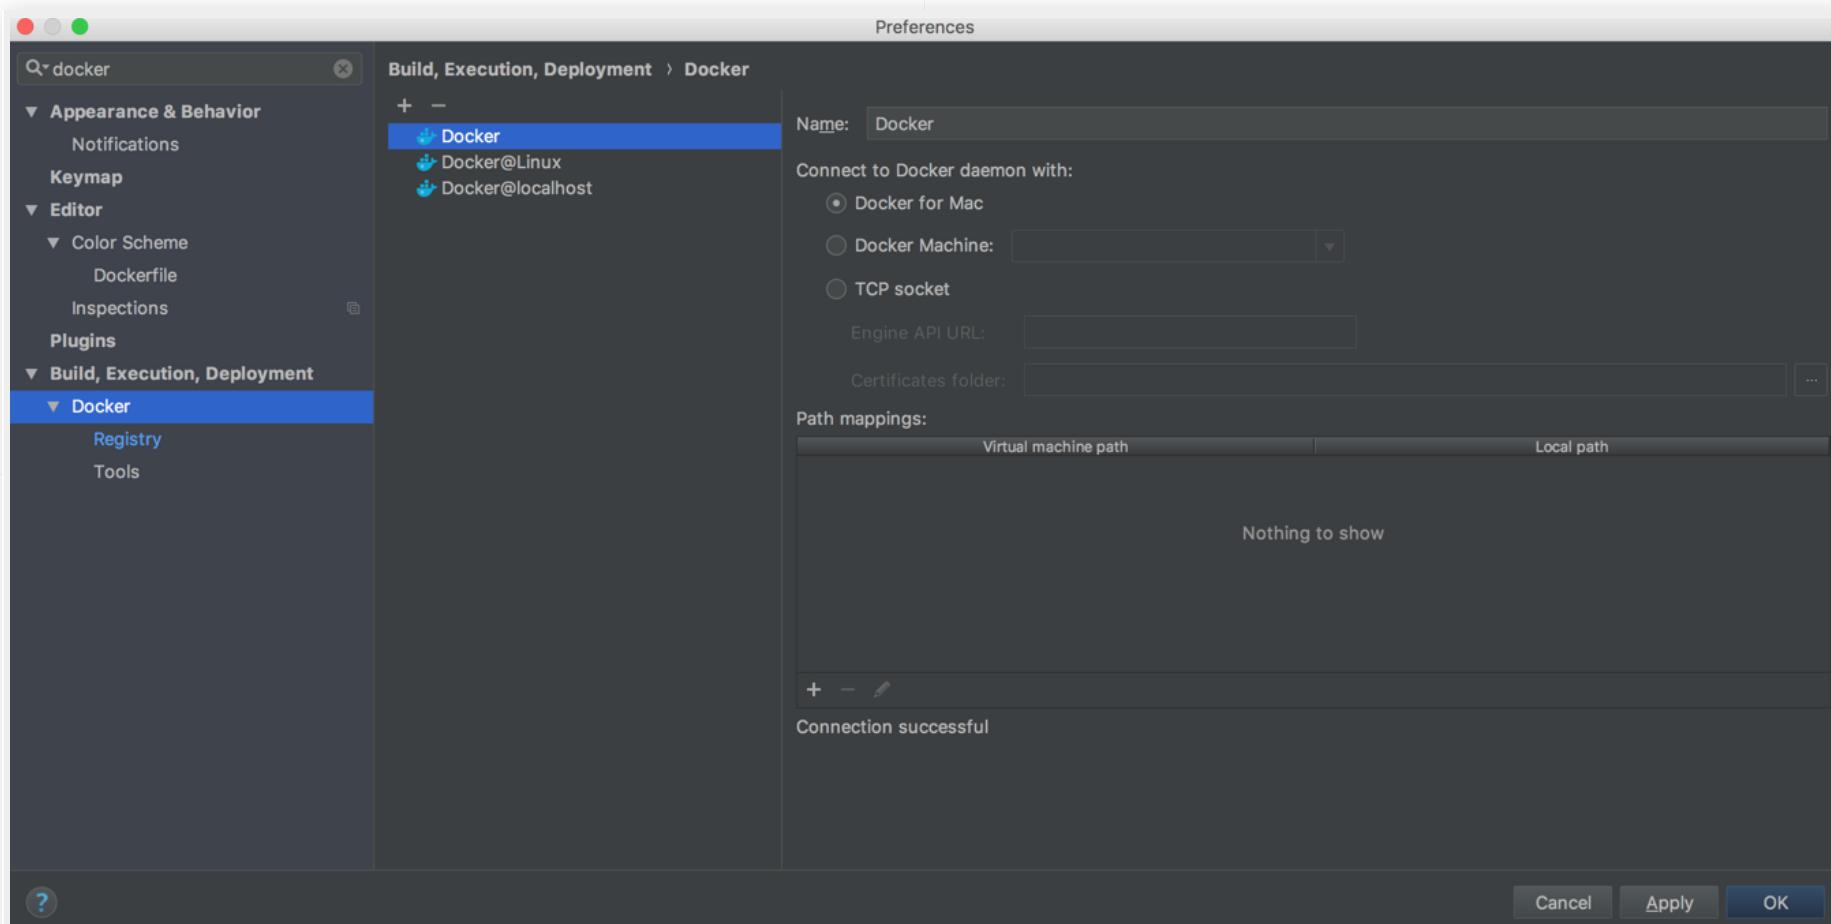
CONFIGURAZIONE SU VSCODE 1/2

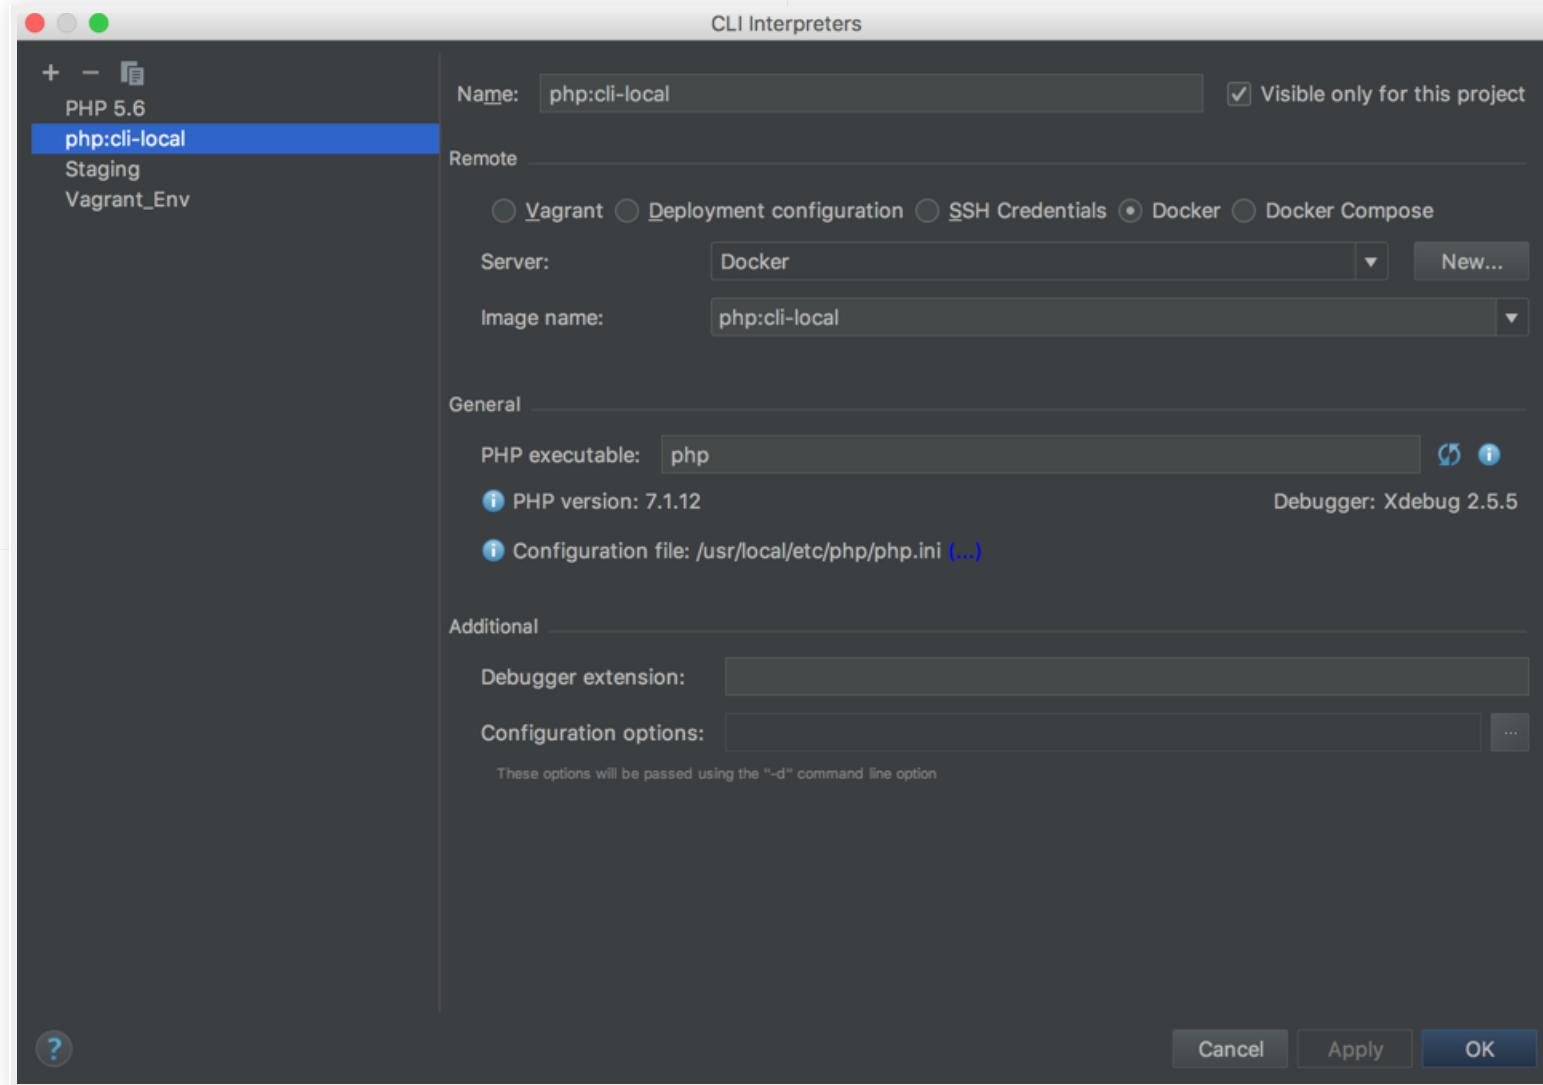
```
[  
  "php.validate.enable": true,  
  "php.validate.executablePath": "${rootPath}/.vscode/bin/php.sh",  
  "php.validate.run": "onSave",  
  "php.suggest.basic": false  
}  
 Giuliano Latini, 3 months ago • Create Workspace Settings for PHP
```

CONFIGURAZIONE SU VSCODE 2/2

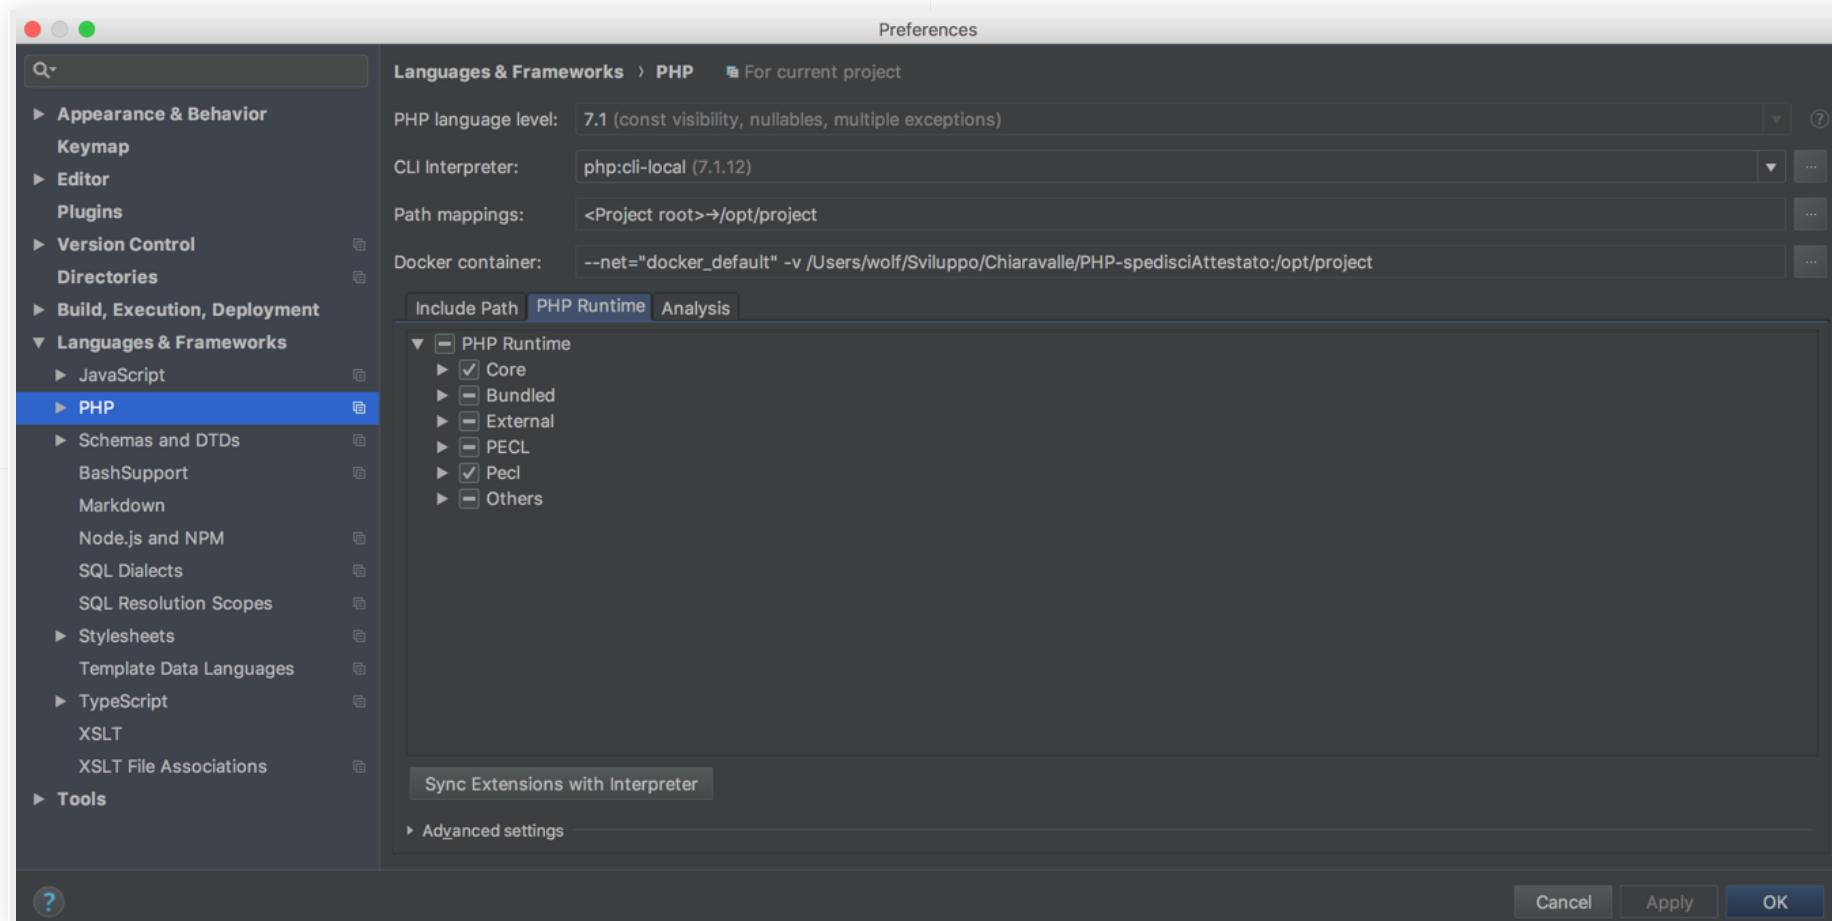
```
#!/bin/sh
# This is a comment!
_heads=${@:1:$((#$ - 1))}
_tail=${@:$#}
_tail=${_tail##$(pwd)}
_tail=${_tail#\//}
/usr/local/bin/docker run -i --name php_cli \
--rm -v $(pwd):/app -w /app php:cli-local php $_heads $_tail
```

CONFIGURAZIONE SU PHPSTORM 1/3





CONFIGURAZIONE SU PHPSTORM 3/3



ask who? discover
where? how? why? challenge who?

QUESTIONS

ask who? discover
when? knowing clues why? known
investigation ? ask
what?

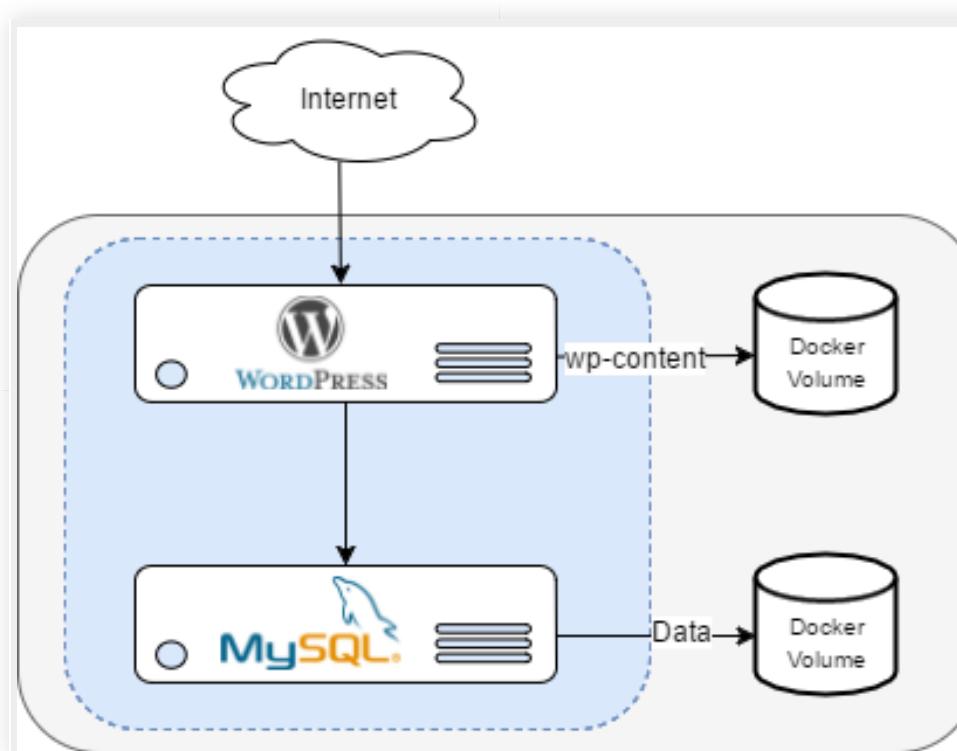
GESTIONE DI SISTEMI

ESEMPIO DI UN'INFRASTRUTTURA WORDPRESS ORGANIZZATA A SERVIZI

DEFINIZIONE FORMALE IMPLEMENTATA CON DOCKER

WORDPRESS IN DOCKER

SCHEMA GENERALE

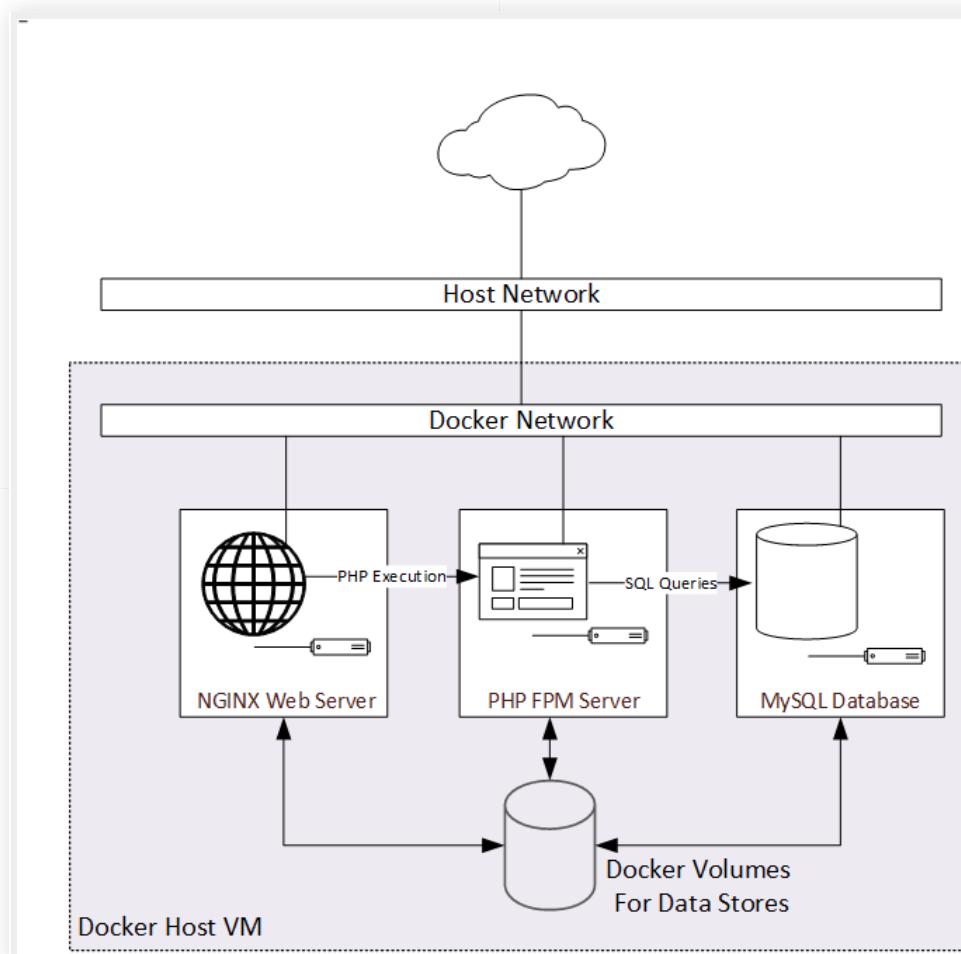


SCHEMA GENERALE (DOCKER-COMPOSE.YML)

```
version: '2'
services:
  wordpress:
    image: wordpress
    links:
      - mysql
    environment:
      - WORDPRESS_DB_PASSWORD=secretpassword
      - VIRTUAL_HOST=example.com
    expose:
      - 80
  mysql:
    image: mysql:5.7
    environment:
      - MYSQL_ROOT_PASSWORD=secretpassword
```

From: ./06_GestioneSistemi/code/01/docker-compose.yml

SCHEMA IMPLEMENTATIVO

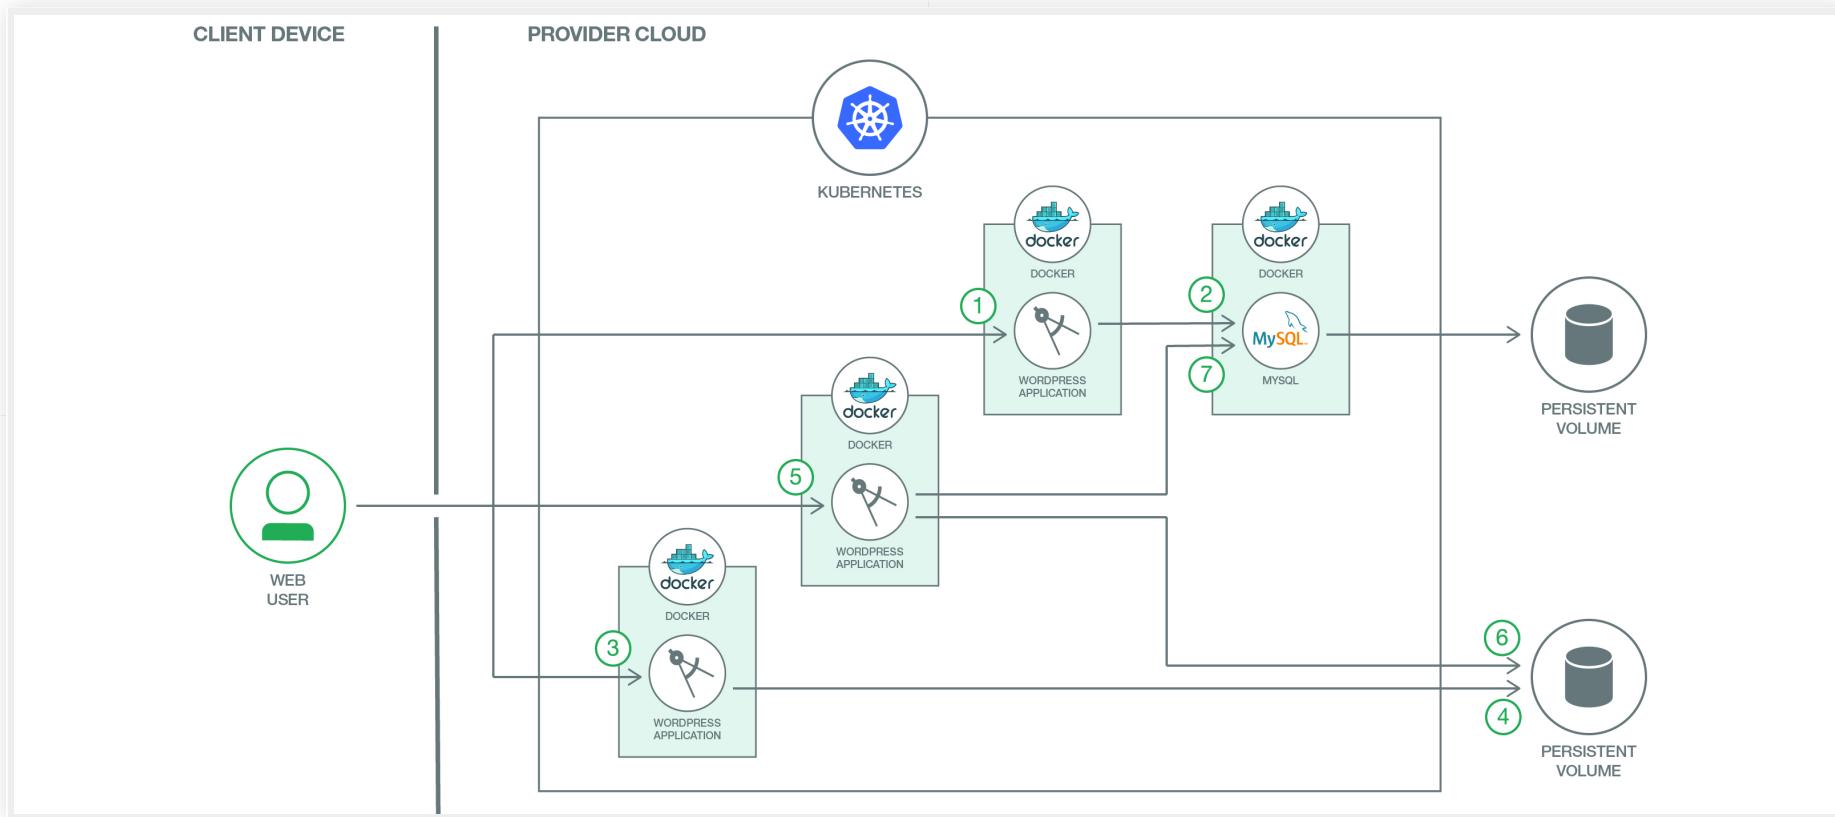


SCHEMA IMPLEMENTATIVO (DOCKER-COMPOSE.YML)

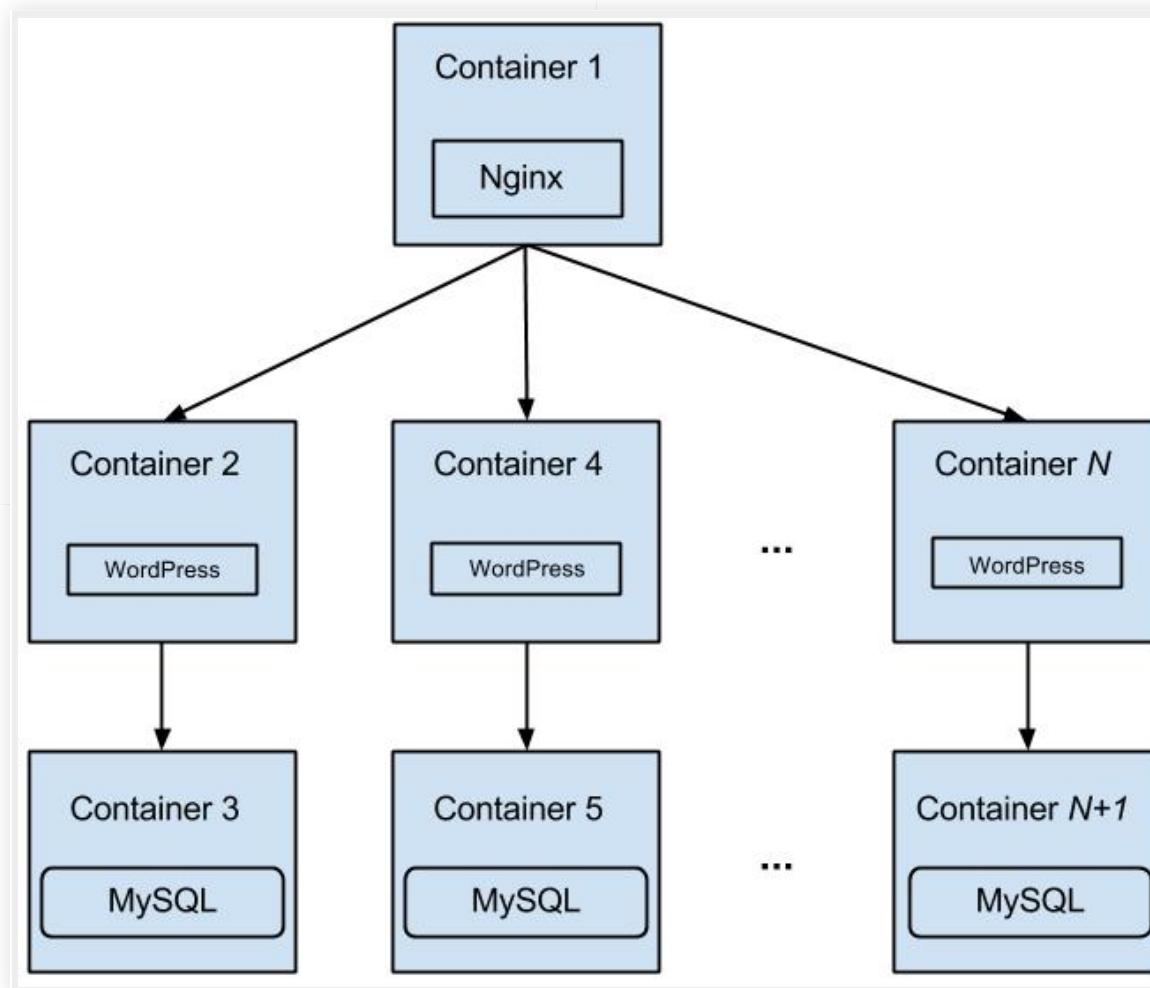
```
version: '3'
services:
  web:
    image: nginx:alpine
    volumes:
      - "./etc/nginx/default.conf:/etc/nginx/conf.d/default.conf"
      - "./etc/ssl:/etc/ssl"
      - "./web:/var/www/html"
      - "./etc/nginx/default.template.conf:/etc/nginx/conf.d/default.template.conf"
    ports:
      - "8000:80"
      - "3000:443"
    environment:
      - NGINX_HOST=${NGINX_HOST}
    command: /bin/sh -c "envsubst '$$NGINX_HOST' < /etc/nginx/conf.d/default.template.conf"
```

From: ./06_GestioneSistemi/code/02/docker-compose.yml

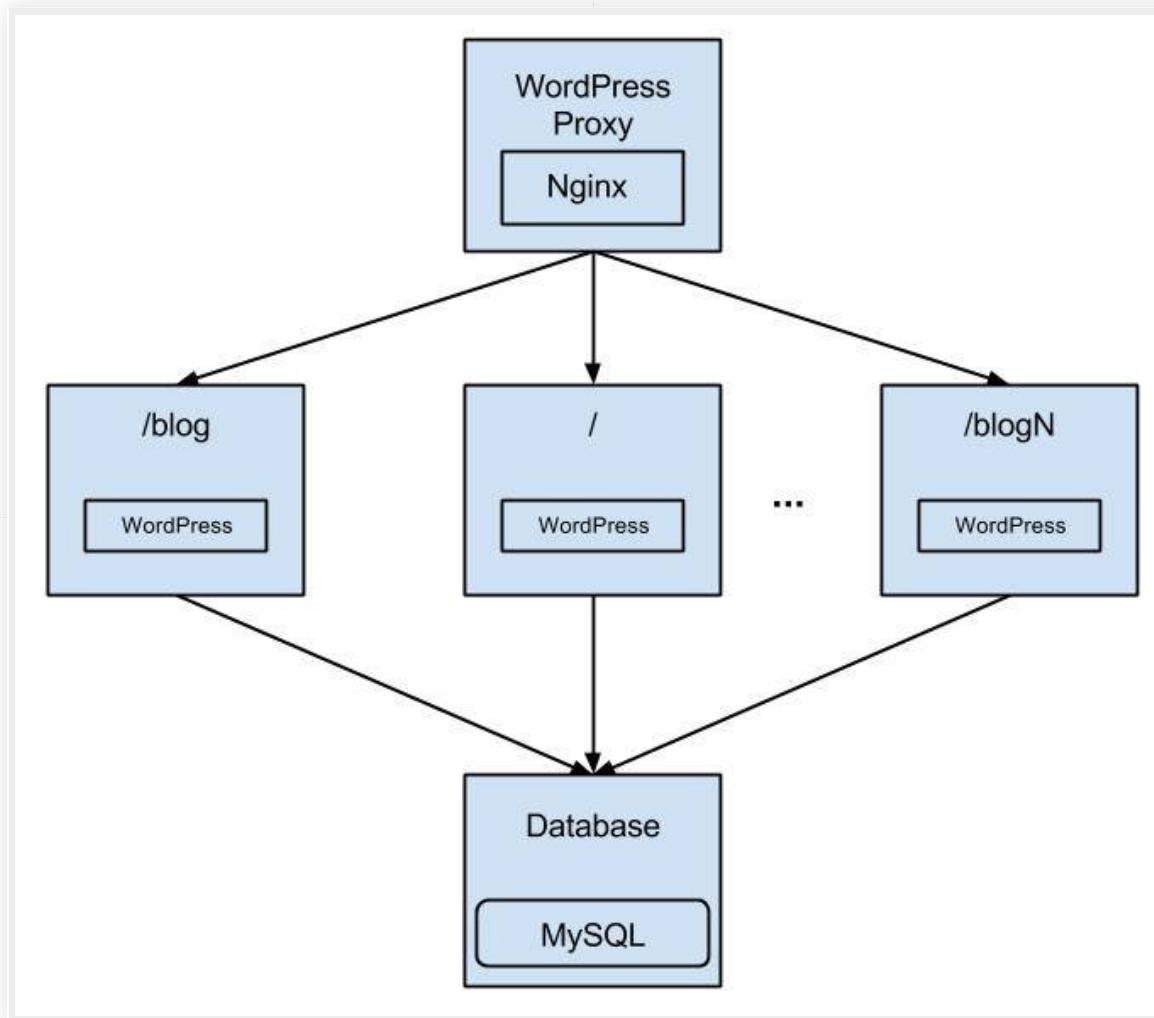
ARCHITETTURA SCALABILE IN DINAMICO



PATTERN ARCHITETTURALE MULTIDBMS



PATTERN ARCHITETTURALE MONODBMS



ask who? discover
where? how? why? challenge who?

QUESTIONS

ask who? discover
when? knowing clues why? known
investigation ? ask
what?

DOCKER SWARM

**LA CLUSTERIZZAZIONE ED ALTA
AFFIDABILITÀ**

**COME IL SISTEMA RISPONDE E SI ADATTA ALLE
SOLLECITAZIONI ESTERNE IN UN SISTEMA DI CALCOLO**

SERVIZIO CLOUDGARR

IL CLOUD PER LA COMUNITÀ GARR

GARR Federated Cloud

Giuseppe Attardi
giuseppe.attardi@garr.it

Dipartimento CSD
Consortium GARR

Demo



GARR Federated Cloud

The **GARR** Federated Cloud offers cloud services to the Italian academic and research community. **GARR** coordinates a federation of clouds, located in national datacenters owned by members of the GARR community, which participate to the federation by sharing resources and services.

The **GARR** cloud allows creating and managing Virtual Machines (**IaaS**) as well as stacks and applications (**PaaS**).

[Go to the dashboard](#)

Virtual Machines

The **GARR** Cloud delivers **virtual machines** running in the data centers of the *GARR Federated Cloud* connected through the **GARR** high speed fiber network. The **GARR** Cloud provides tools for self provisioning computing resources and deploying applications and services, enabling scaling from single instances to clusters of integrated and load-balanced cloud computing.



Register to GARR Federated Cloud



Create a new account

Federated Cloud

Objectives

- Facilitate **transition towards cloud computing**
- Allow **resource sharing**, maintaining **control of use**
- Exchange **best practices** on management and use
- Evolve towards **native cloud applications**
- Expand **catalogue of cloud applications**

GARR Commitments

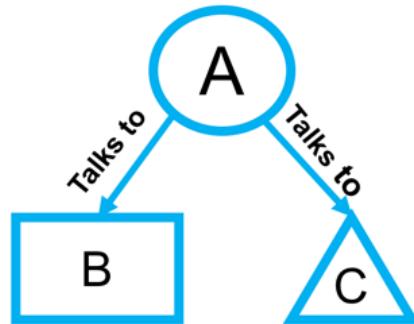
- Architecture Design
- Ready to use OpenStack Distro
- OpenSource Code Base (git.garr.it)
- Upgrades and Maintenance
- Solution for multiple tenancy
- Federation and Delegation
- Federated Authentication
- Asset Management

TCO: after ownership

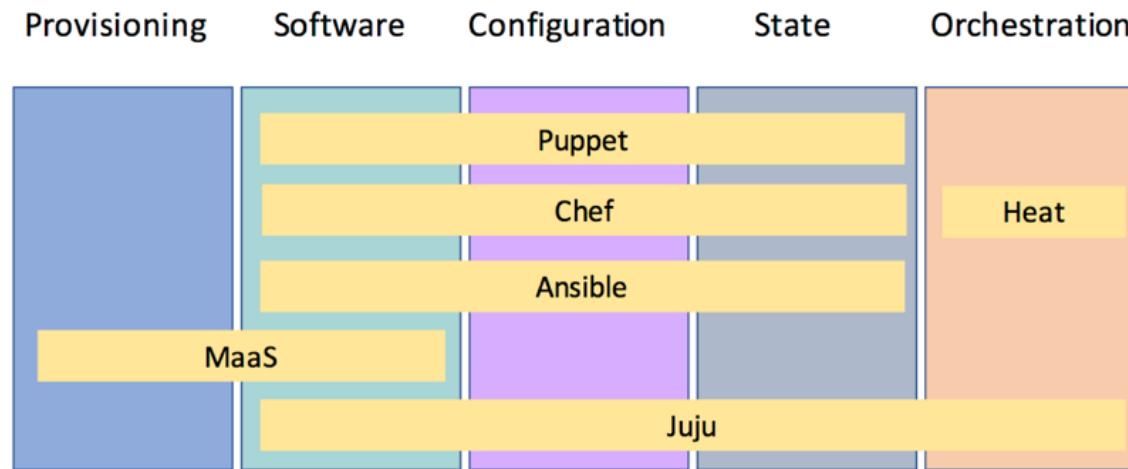
- If you buy services, at the end **what does it remain to you?**
 - Just the experience on using somebody else's product
- If you build your own infrastructure, at the end what remains to you?
 - The infrastructure
 - The experience in building the infrastructure
 - The experience on using the infrastructure

Declarative Modeling

- App A requires:
 - X GB memory and Y CPU
 - N GB storage
 - talking with B and C
 - An URL endpoint
 - To run locally, close to B

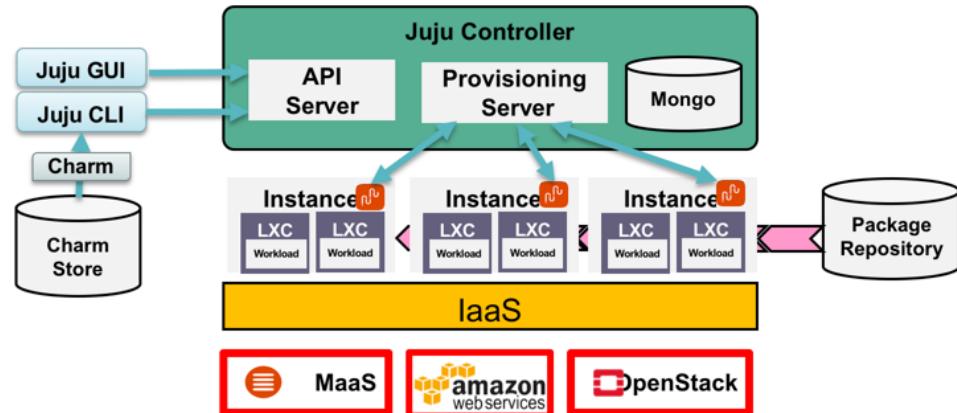


Automation Tools

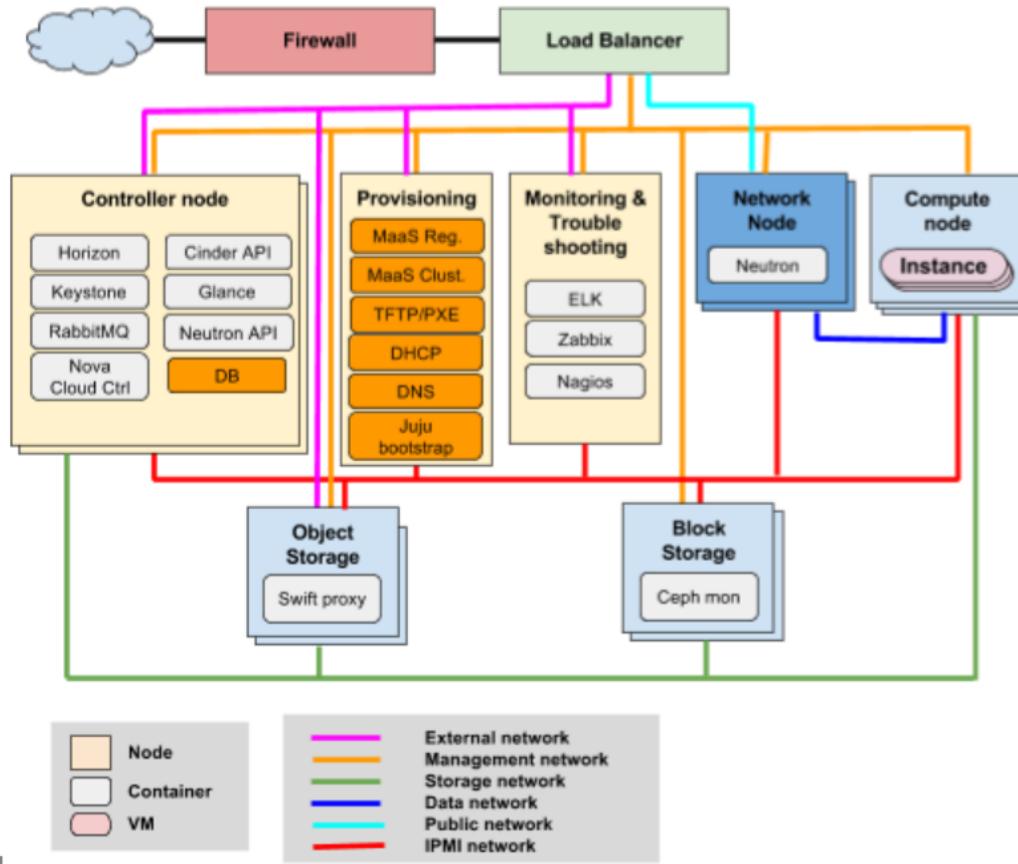


Juju Architecture and Workflow

- **Reactive engine**, triggered by events that fire corresponding handlers
- **Multiple handlers** may match and will be run in an undetermined order
- Handlers may cause additional events
- The engine **runs until convergence** to a stable state



- Widely used and well supported Cloud Computing software
- Over 45.000 developers world wide
- Complex to manage
- Designed a Reference Architecture:
 - Declarative modeling
 - Easy to configure and to replicate
 - Managed with automated orchestration tools



Status

● Resources

- ~9000 vCPU
- 10 PB Storage

● Usage

- Over 700 users
- Over 1000 VM

● Guarantees

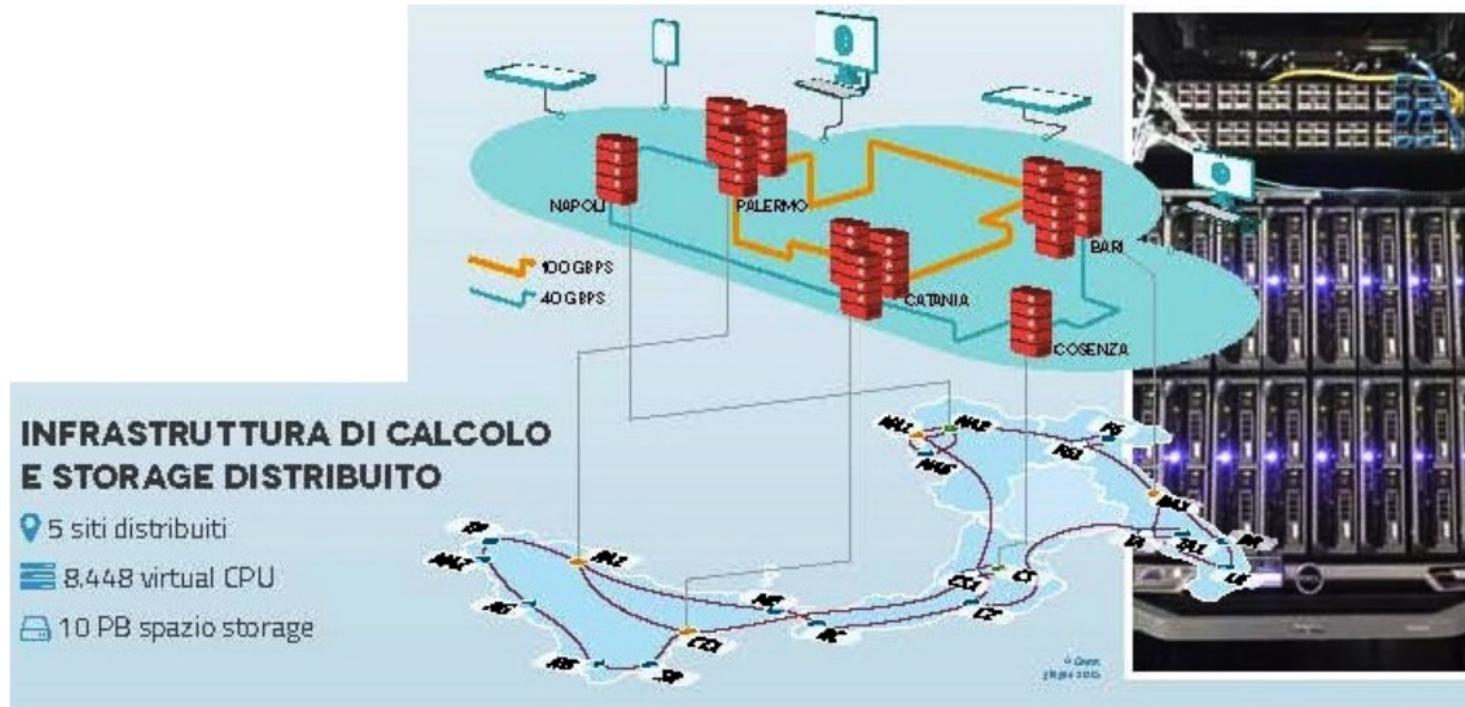
- Service Continuity
- Data Protection

Usage

[Download CSV Summary \(?start=2016-11-01&end=2017-03-22&format=csv\)](#)

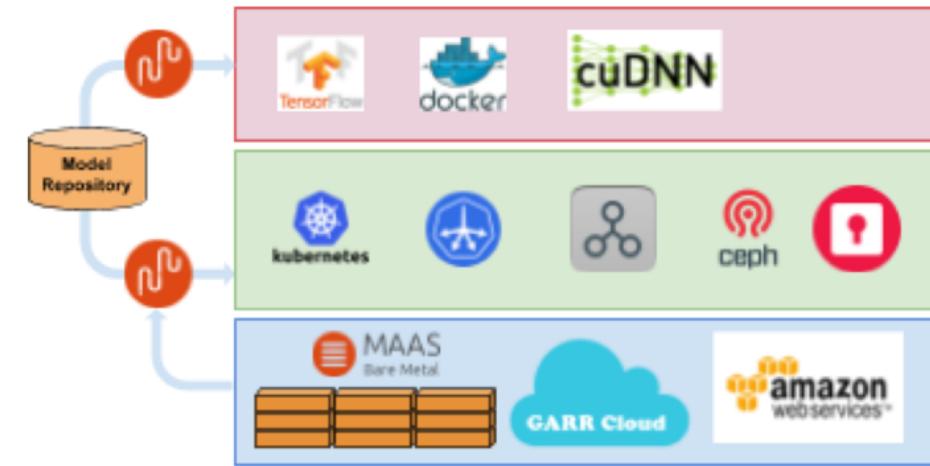
| Project Name | VCPUs | Disk | RAM | VCPU Hours ⓘ | Disk GB Hours ⓘ | Memory MB Hours ⓘ |
|--------------|-------|-------|--------|--------------|-----------------|-------------------|
| isti | 368 | 2.8TB | 656GB | 164333.47 | 1322184.30 | 279763673.66 |
| unipa | 336 | 7.5TB | 1.2TB | 693206.14 | 16207815.52 | 2692199747.8 |
| ws2017ipv6 | 128 | 2.5TB | 256GB | 14955.82 | 299116.37 | 30629516.41 |
| lns-prj1 | 54 | 543GB | 71.5GB | 118279.27 | 1629777.45 | 189620873.70 |
| garrdemo318 | 32 | 20GB | 32GB | 16914.14 | 10571.34 | 17320083.68 |
| INFN-FI | 32 | 40GB | 16GB | 601.92 | 771.89 | 309779.80 |
| wsosadmin | 27 | 540GB | 54GB | 4510.02 | 90200.37 | 9236517.59 |
| GEANT | 25 | 481GB | 48.5GB | 29681.28 | 591502.48 | 60615623.66 |
| infn-vlabs | 24 | 480GB | 48GB | 8052.53 | 161050.61 | 16491582.24 |
| demo | 17 | 340GB | 34GB | 9877.99 | 197559.85 | 20230128.54 |
| SSSUP | 16 | 80GB | 16GB | 12651.09 | 63258.04 | 12954892.40 |
| garrdemo125 | 8 | 100GB | 128GB | 13429.14 | 167864.31 | 220023110.77 |

Deployment Infrastructure



Container Platform for AI

- Automated deployment on bare metal, AWS or other clouds by Juju
- Workloads deployed by Juju
- Distributed storage system using Ceph
- NFS cluster for sharing big data
- Docker containers managed Kubernetes

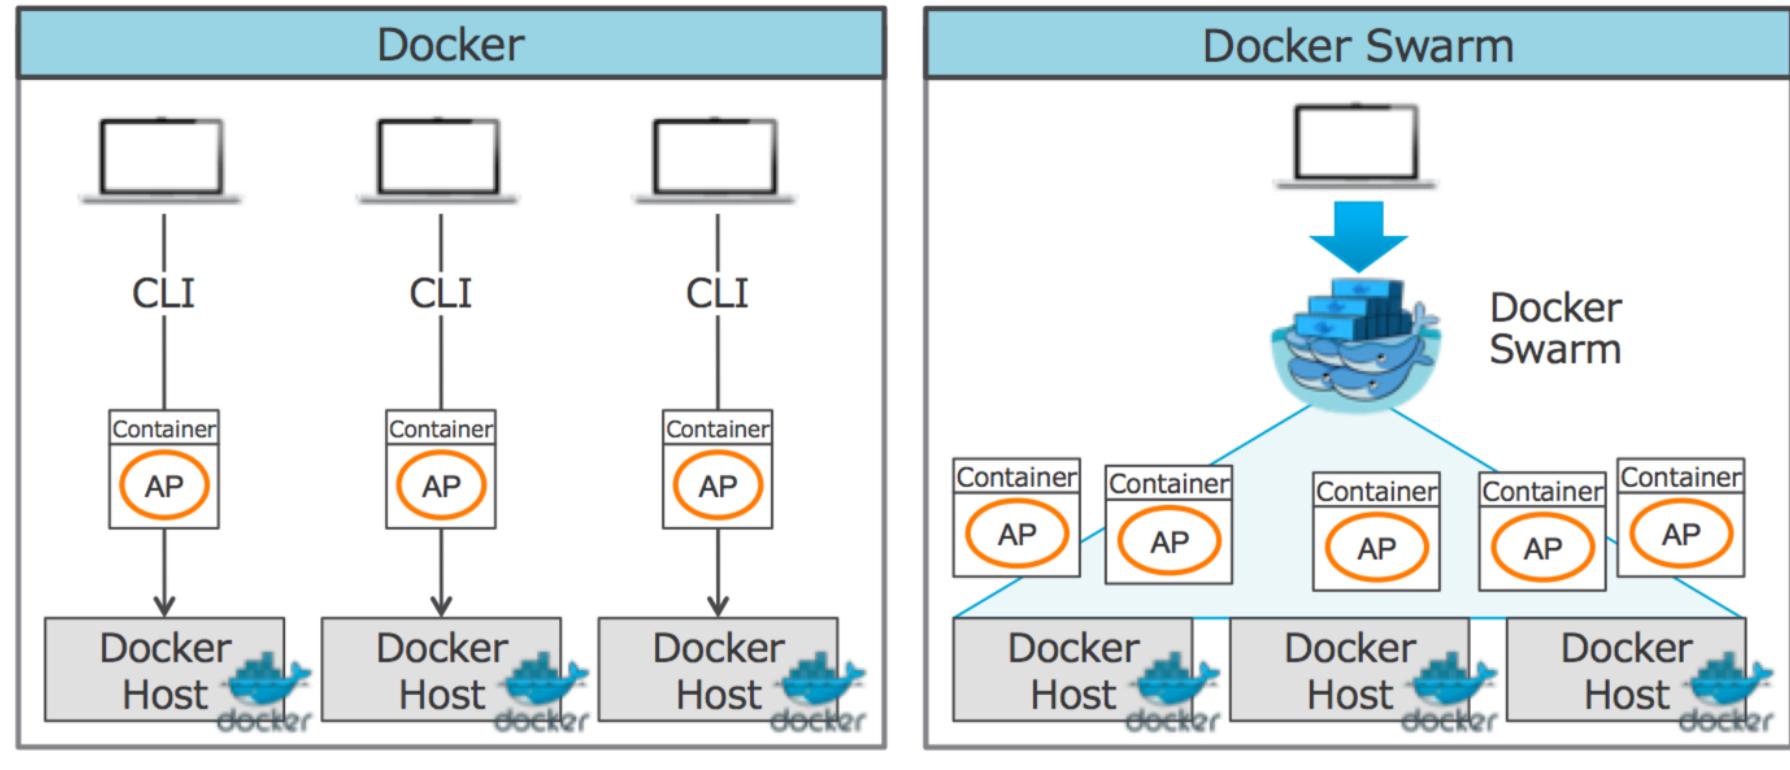


- Build a community of users and developers:
 - <https://cloud.garr.it/community/>
- Build a shared Catalogue of services
- Examples:
 - Moodle as a Service
 - Jupyter Notebooks as a Service

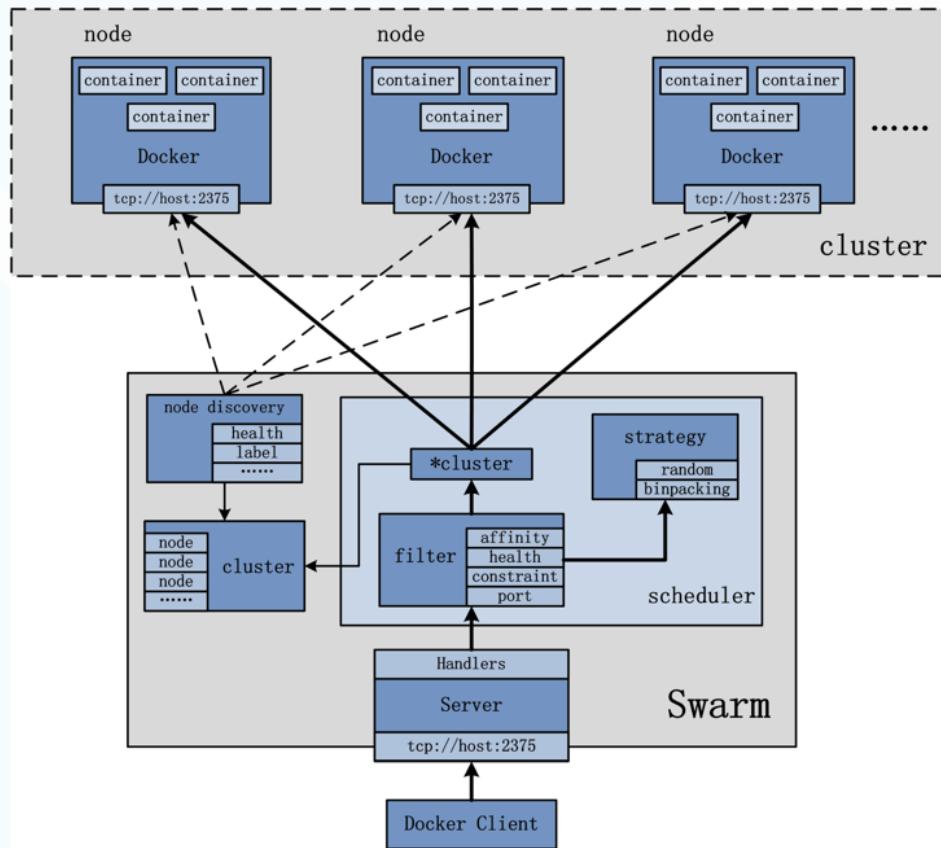
SWARM

CLASTERIZZAZIONE DI RISORSE

Docker Engine vs Docker Swarm

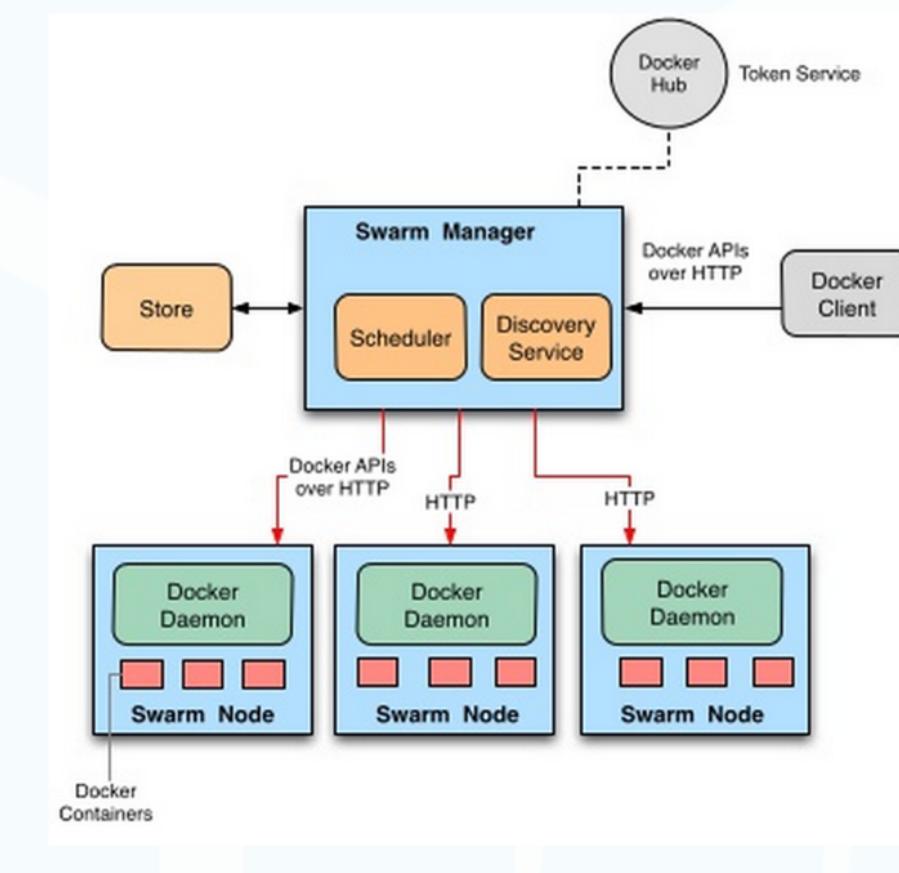


ARCHITETTURA FISICA DOCKER SWARM



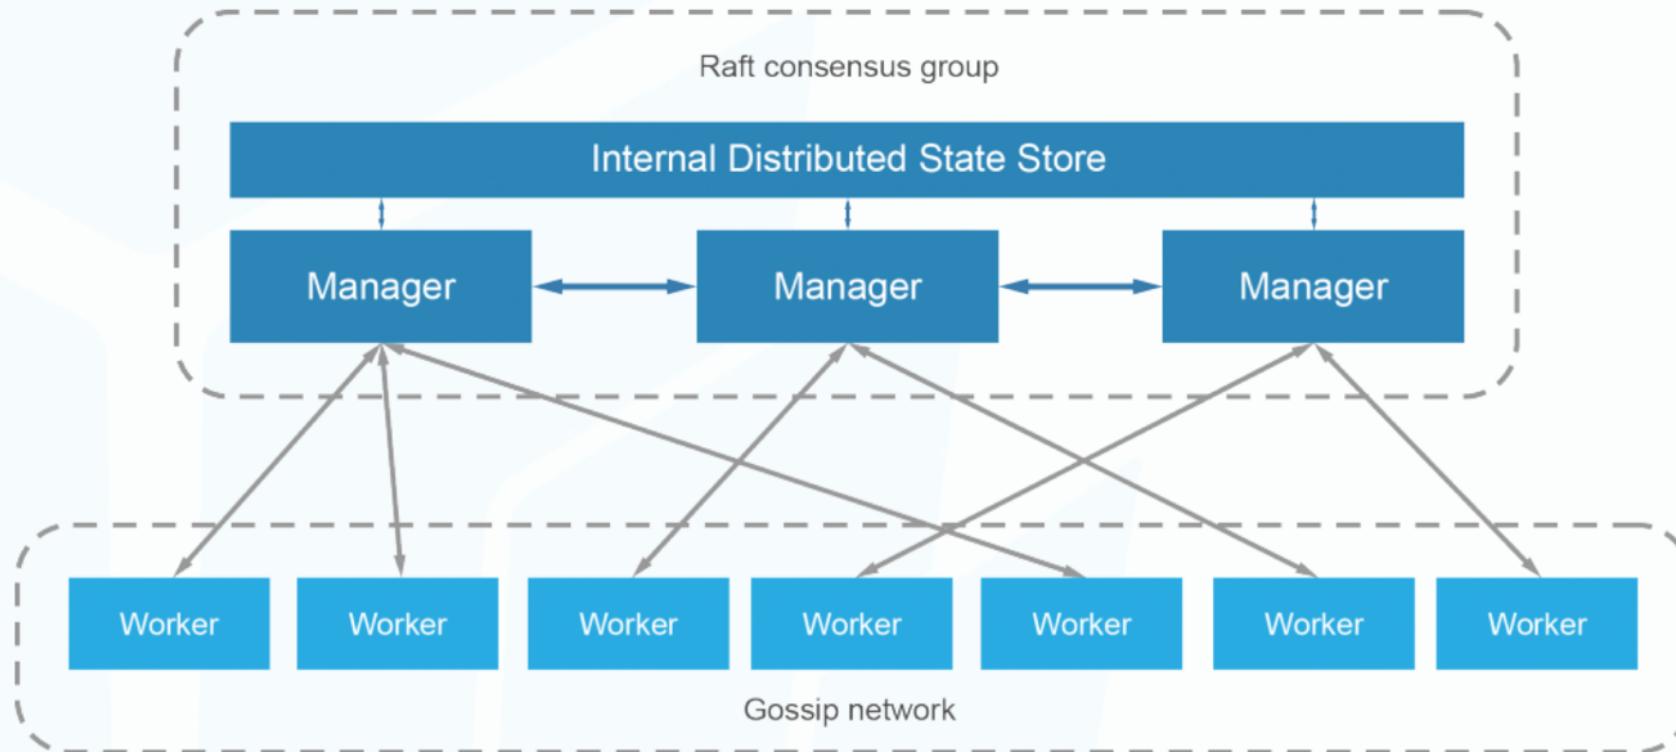
Swarm master organizza una rete di docker-host garantendo raggiungibilità ed accessibilità così da disporre tramite il docker client di un punto di accesso unico a tutte le risorse.

ARCHITETTURA FUNZIONALE DI UN CLUSTER

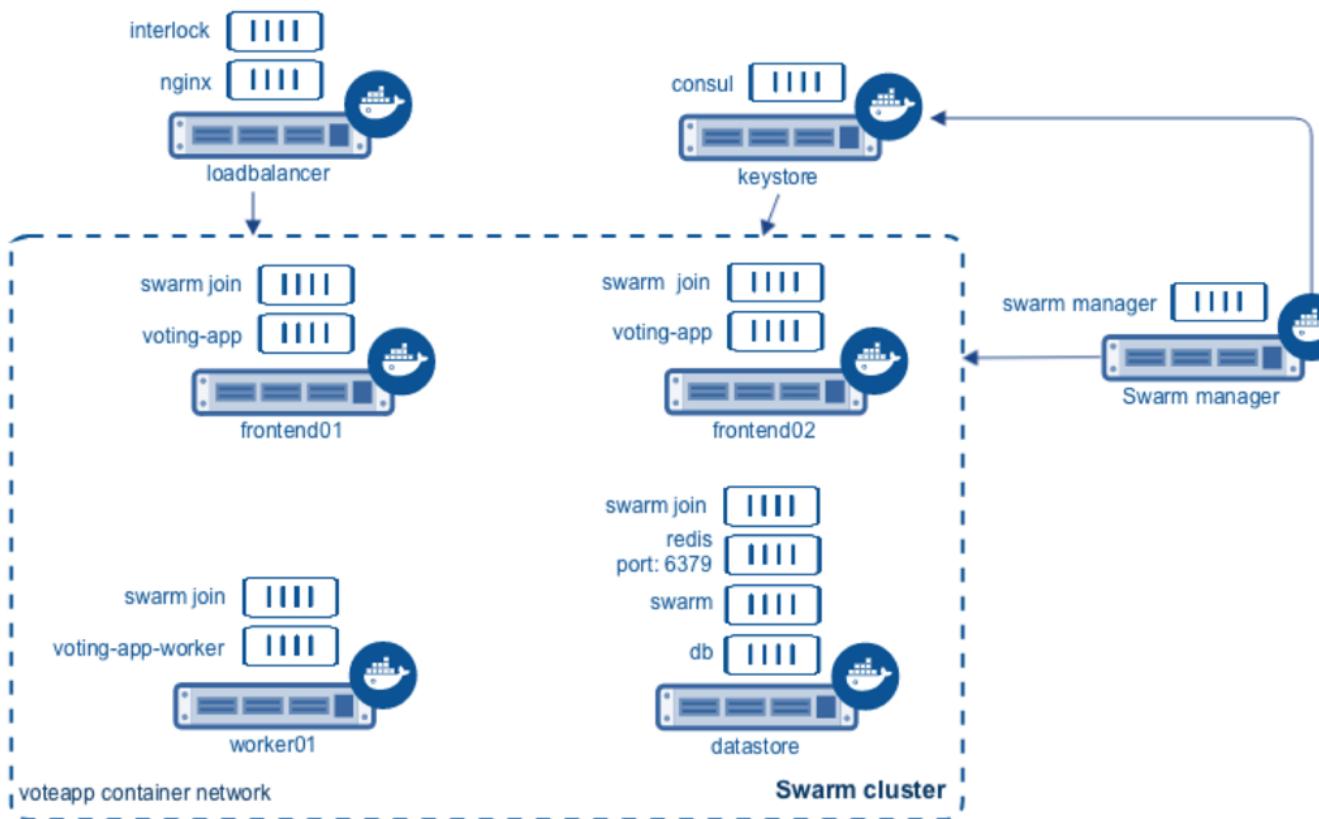


- Docker Client facilità l'accesso alle risorse del cluster.
- Docker Hub è il repository da cui attingere per le immagini da eseguire come container dal Docker Daemon.

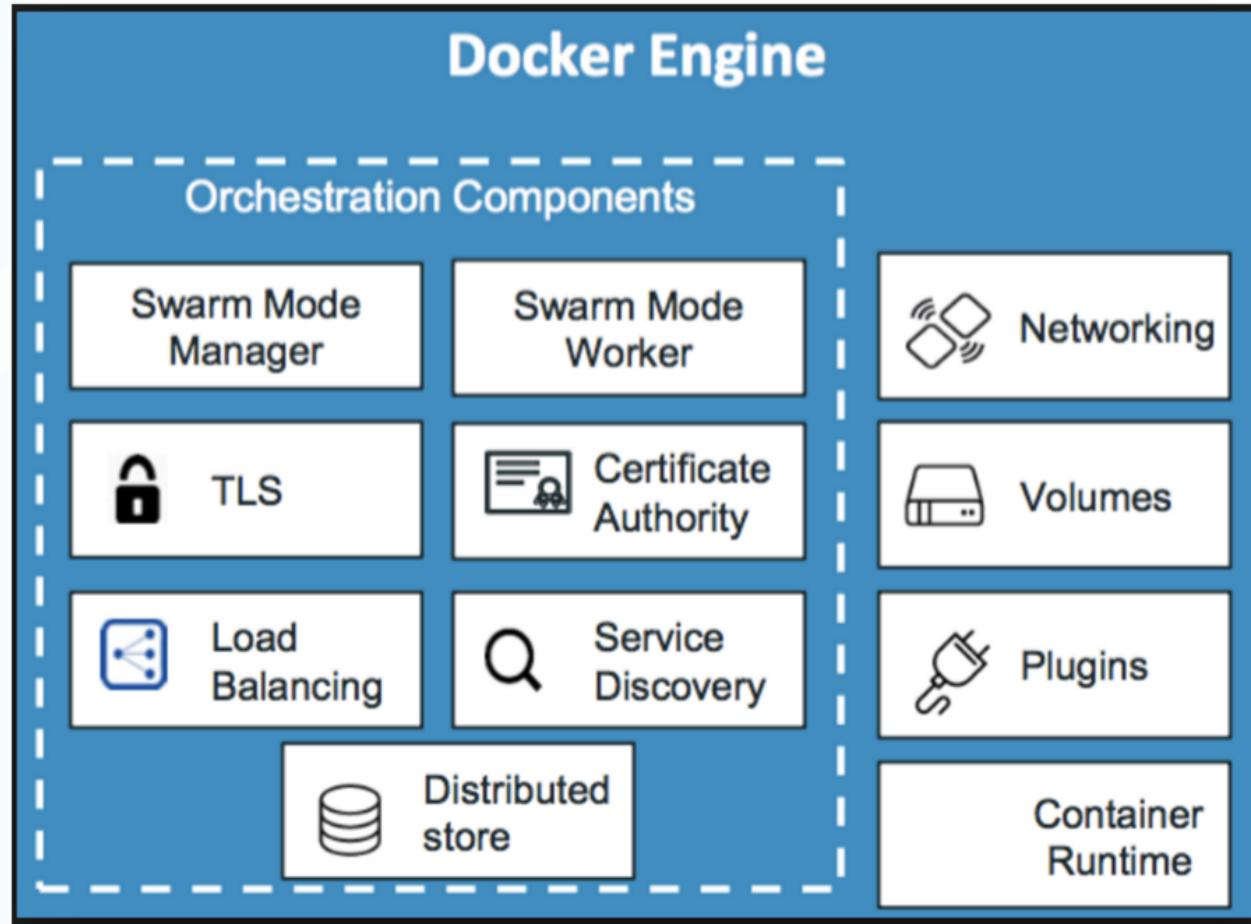
ARCHITETTURA FUNZIONALE DI UN CLUSTER

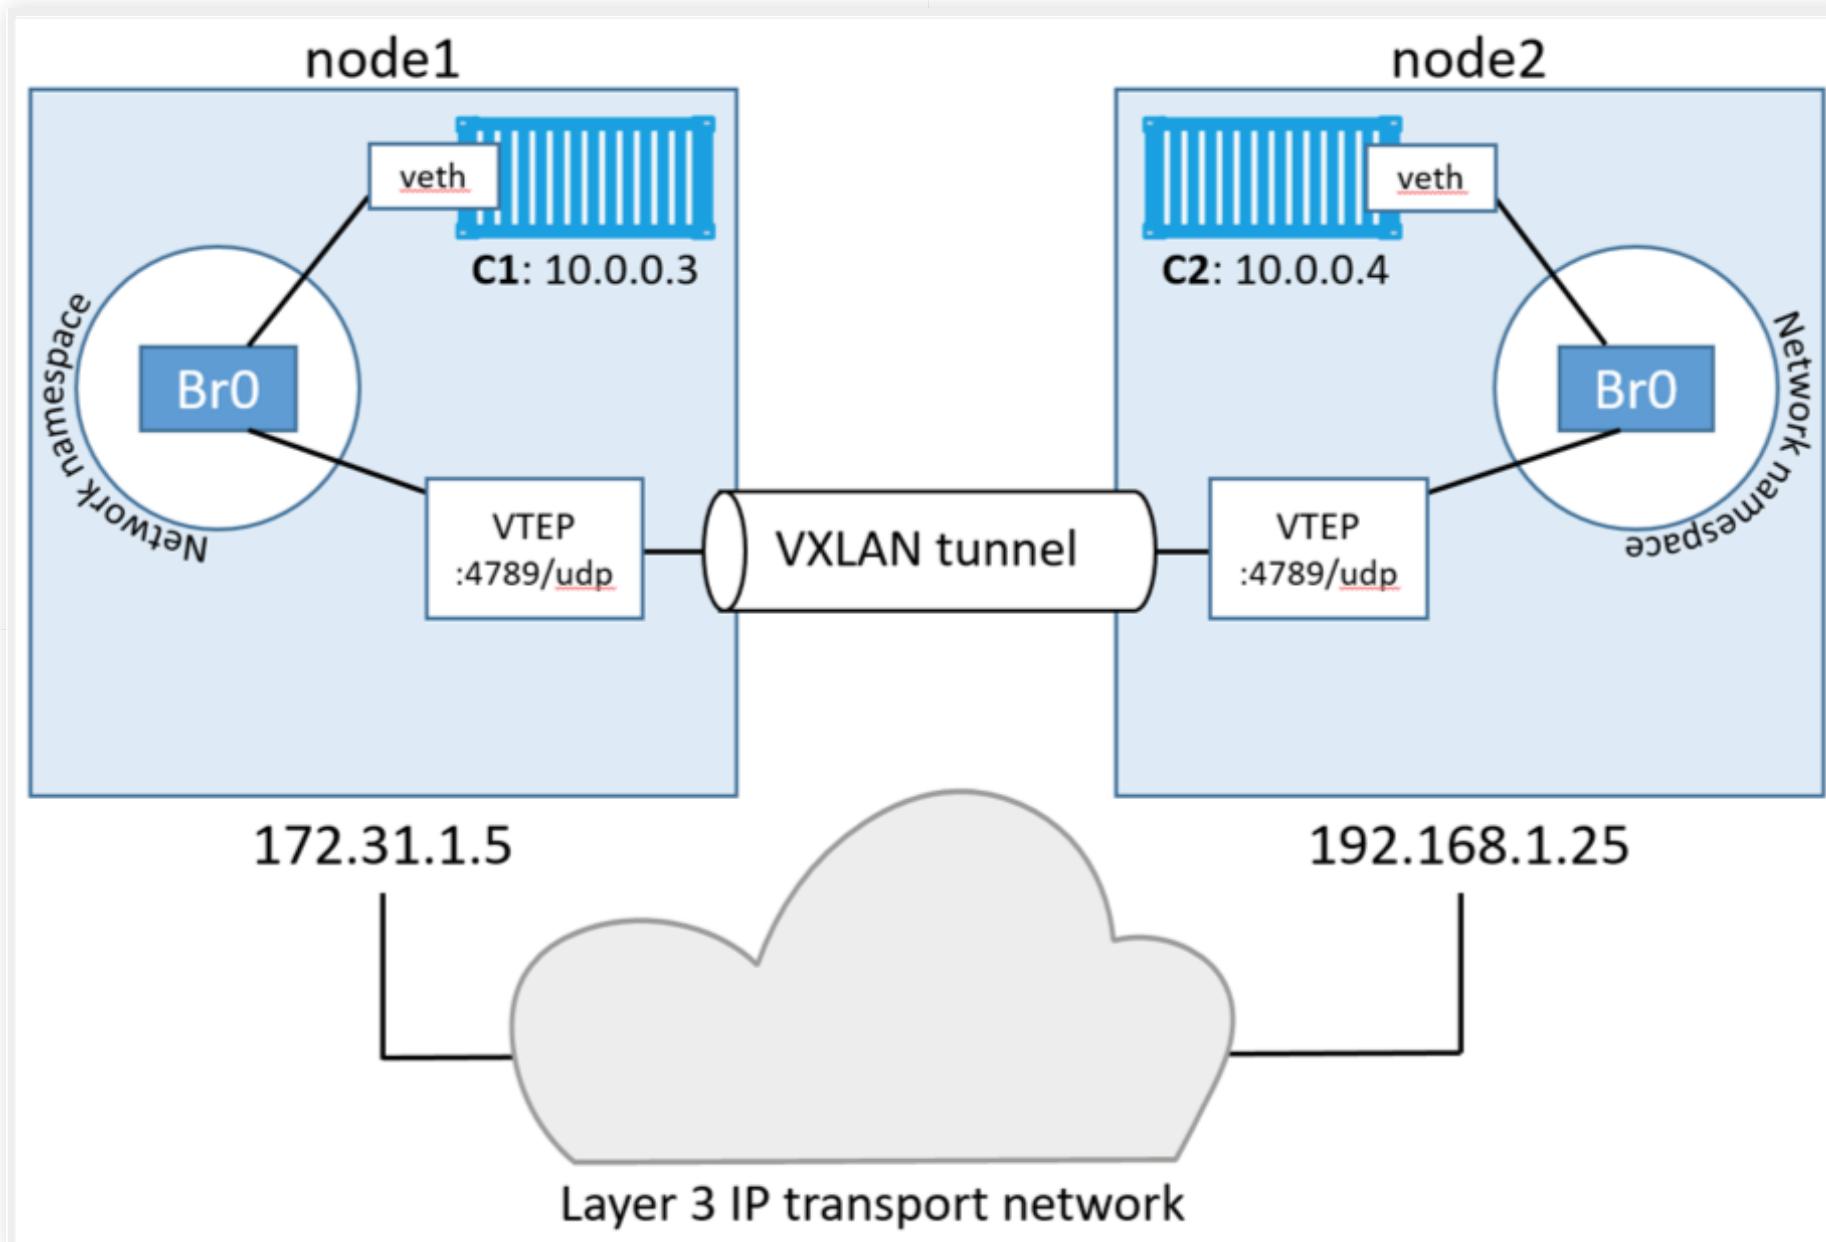


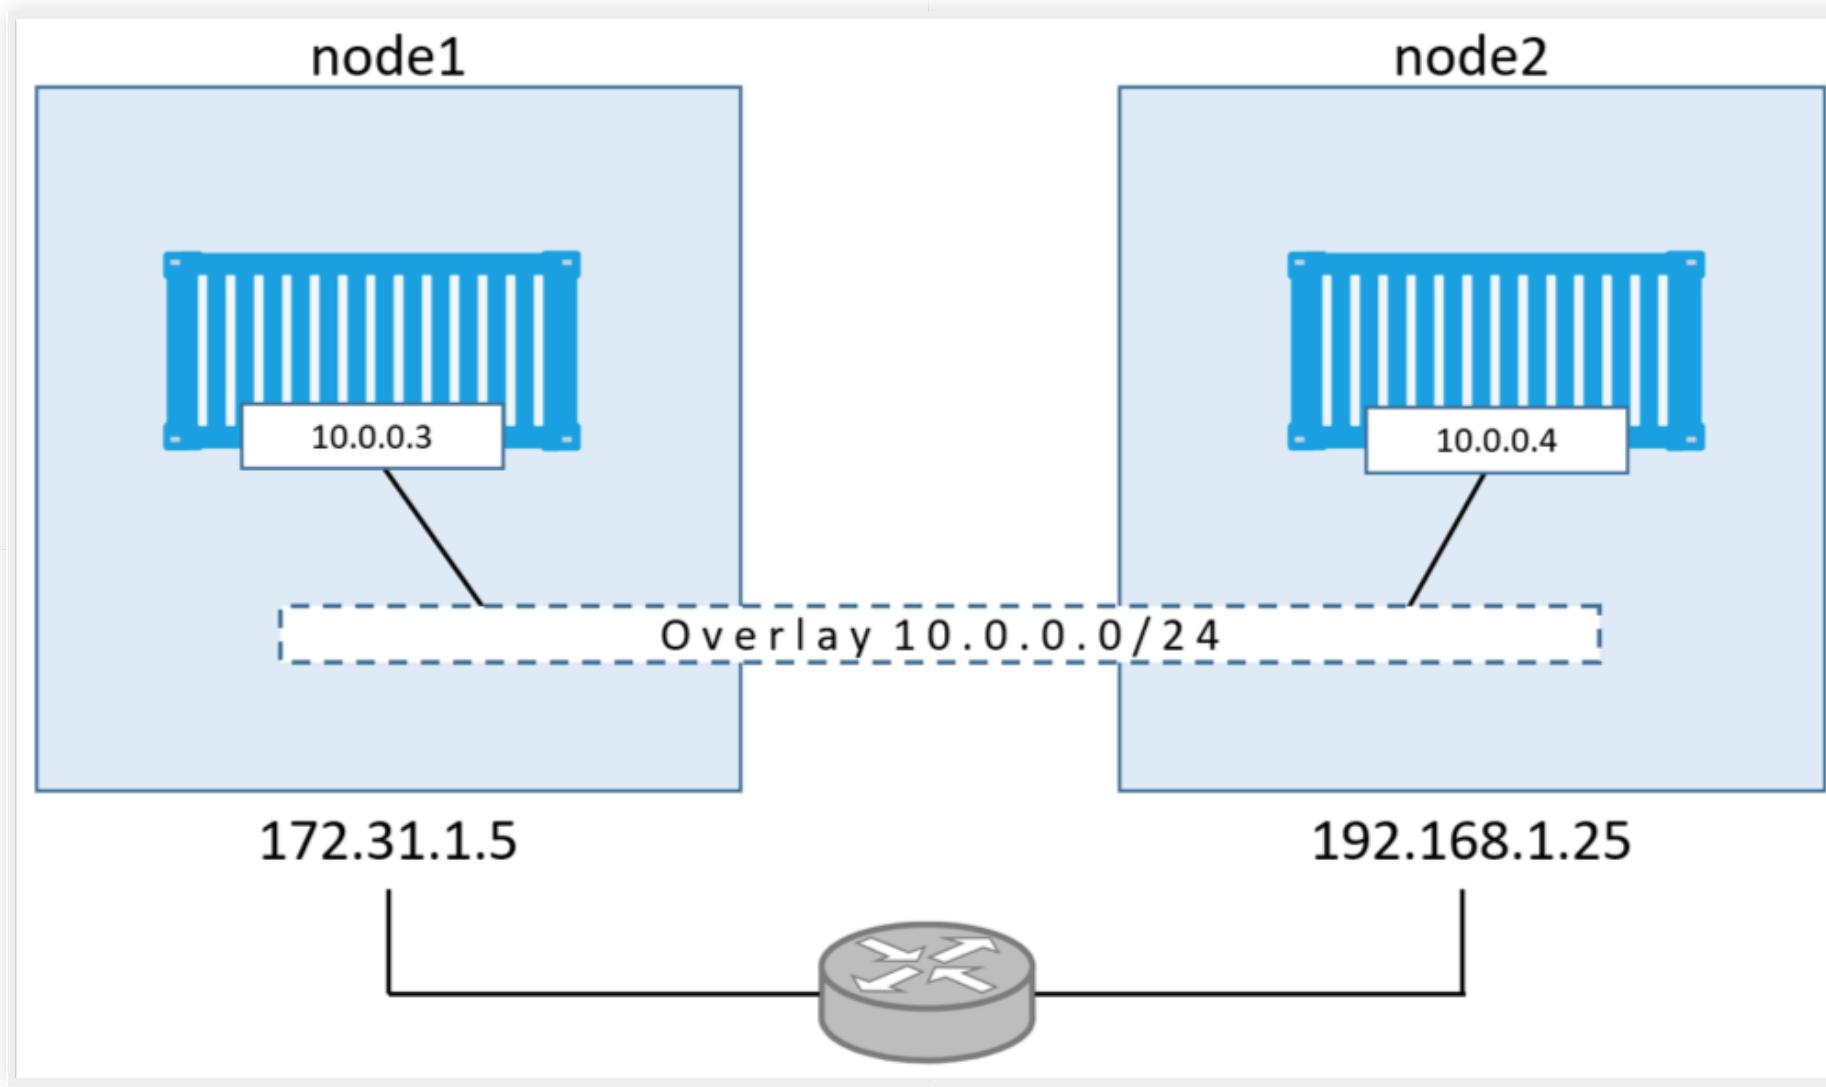
ARCHITETTURA FUNZIONALE DI UN CLUSTER



LA MAGIA IN DOCKER ENGINE







ask who? discover
where? how? why? challenge who?

QUESTIONS

ask who? discover
when? knowing clues
investigation ? why? ask
what known investigation